

Drug Lords in STEM

Maria Montoya-Aguirre*

September 19, 2026

[Click here](#) for the latest version

Abstract

Education is widely viewed as a protective factor against criminal involvement because it improves legal labor-market opportunities. But what happens when illegal markets demand skilled labor? I study how a boom in chemistry-intensive illegal drug production reshaped human capital investment in Mexico. China's 2019 ban on fentanyl pushed its synthesis into clandestine Mexican laboratories, expanding synthetic drug production and the demand for chemistry skills. Using a difference-in-differences design, I compare commuting zones with pre-existing synthetic-drug infrastructure to those without before and after the ban. Water-quality monitoring data confirm that exposed areas exhibit the byproducts of drug synthesis. Exposure reduces school dropout, increases time spent studying, and raises the share of students who start the final year of high school. Tertiary enrollment in chemistry-related STEM programs also rises, with no change in other fields, confirming a skill-specific shock rather than an income effect. The demand for chemistry skills spills over to the legal sector: exposure increases the wage premium on chemistry-related degrees and employment in chemistry occupations. Meanwhile, older credentialed chemists disappear from the data, consistent with entry into concealed production.

JEL Codes: I26, J24, K42, O15

Keywords: human capital, education, organized crime, drug trafficking, skill-biased technical change

*Paris School of Economics and Université Paris 1 Panthéon-Sorbonne maria.montoya@psemail.eu. I am very grateful for the generous guidance and support from David Margolis, Oliver vanden Eynde, Nina Guyon and Maria Micaela Sviatschi. Thank you to seminar participants at CEPR Development Economics Symposium, Chicago/LSE Conference on the Economics of Crime and Justice, Paris School of Economics, Texas Economics of Crime Workshop, WEC Jr. at the Institute for Fiscal Studies, ViCE seminar, and EEA-ESEM conference. This work benefited from the EUR grant ANR-17-EURE-0001 and a Horowitz Foundation Dissertation Award.

1 Introduction

Education is widely viewed as a protective factor against criminal involvement. The canonical framework in the economics of crime, originating with [Becker \(1968\)](#), holds that education raises the opportunity cost of crime by improving legal labor market prospects, and a large body of empirical work supports this relationship ([Lochner and Moretti, 2004](#); [Machin et al., 2011](#); [Machin and Sandi, 2025](#)). Yet the framework rests on an implicit assumption: that criminal activity primarily demands low-skilled labor, for which education confers no productive advantage. This fits street crime and much of the physical work of the drug trade—cultivating crops, moving contraband, providing muscle for territorial control. But it need not hold when criminal organizations modernize their operations and begin to demand skilled workers. This paper examines the effects of a skill-biased shift in the technology of an illegal market on the human-capital decisions of the youth exposed to it.

I study a setting where technological change in an illegal market clearly raised the demand for a specific skill. Over the past two decades, Mexican trafficking organizations have shifted from plant-based drugs—marijuana, heroin from opium poppy—toward synthetic drugs, particularly methamphetamine and fentanyl. This shift transforms the skills that criminal production rewards: growing cannabis or poppy requires land, agricultural knowledge, and physical labor, whereas synthesizing methamphetamine or fentanyl requires laboratory equipment, chemical expertise, and the ability to follow—or invent—complex production protocols. This demand for chemistry skill is not merely theoretical. Leaked documents from Mexico’s military intelligence reveal that trafficking organizations actively recruit chemistry professors from universities, and chemistry students and graduates report recruitment attempts ([Garcia, 2023](#); [Dudley, 2023](#)).

To identify the effect of exposure to high-skill criminal opportunities, I exploit a policy shock that accelerated the transition towards illegal synthetic drug production. On April 1, 2019, China implemented a nationwide ban on fentanyl and its analogues, cutting off the finished product that Chinese firms had been shipping directly to American consumers. Demand did not disappear; so, production move elsewhere. Mexican organizations, already operating methamphetamine laboratories, started fentanyl synthesis in-house, expanding synthetic drug production where they had the necessary infrastructure. Using a difference-in-differences design, I compare commuting zones that had clandestine methamphetamine-laboratory seizures before the fentanyl crisis—indicating pre-existing synthetic-drug infrastructure—to those that did not, before and after the ban. The identifying logic is that existing infrastructure gave these zones a comparative advantage in pivoting to fentanyl. This is validated by two facts: post-2019 fentanyl and laboratory seizures concentrate disproportionately in the pre-exposed zones, and independent water-quality monitoring shows those zones exhibit the chemical signature of methamphetamine synthesis—elevated mercury and toluene—confirming that the exposure measure captures genuine production rather than just differential policing of drug laboratories.

A simple model of occupational choice organizes the consequences of this shock on human capital accumulation. I treat chemistry skill as priced in a local labor market in which a legal and an illegal, drug-producing sector bid for the same, slow-to-replenish pool of chemists, and in which illegal production potentially demands legal inputs—precursors, equipment, and laboratory and commercial services. A rise in the illegal demand for chemistry skill also raises its return on the legal market because of poaching as well as through potential input linkages between the legal and illegal sectors. Youth with different propensities for illegal work choose education, field, and sector. The model predicts that educational investment rises in exposed areas, specifically in chemistry-intensive fields. Additionally, this framework allows to use the estimated legal wage response to gauge the otherwise-unobservable illegal demand shift and corresponding rise in the illegal chemistry premium. Since the shock meets a short-run supply of chemists that cannot be replenished quickly, even a relatively small illegal sector can move local wages of chemists and schooling.

The analysis draws on several data sources and spans 2014 to 2024, a symmetric window around the April 2019 ban. Individual-level educational and labor-market outcomes—enrollment, time use, wages, occupation, and educational attainment for school-age youth and prime-age adults—come from Mexico’s National Survey of Occupation and Employment (ENOE), the country’s principal labor-force survey. I also measure grade progression up to high school using *Formato 911*, the administrative census of the Mexican school system, and tertiary enrollment comes from administrative data of the National Association of Universities and Higher Education Institutions (ANUIES). I allocate tertiary education enrollment to students’ origin commuting zones using student-migration patterns from the 2010 Population Census and classify programs by chemistry intensity from detailed curriculum analysis of programs at Mexico’s largest public university. Treatment is built from military records of clandestine-laboratory seizures maintained by Mexico’s Secretariat of National Defense (SEDENA); national water-quality monitoring (RENAMECA) is used to validate measures of synthetic drug production. I construct 1,226 commuting zones in order to use a local labor market as the unit of analysis.

The results overturn the conventional wisdom about education and crime. Following China’s ban, exposure to synthetic-drug production *reduced* school dropout. Middle-school dropout fell by 1.7 percentage points (a 29% reduction relative to the control mean of 5.9%) and high-school dropout by 2.1 percentage points (an 8.6% reduction relative to the control mean of 24.5%). Event-study estimates show no visual evidence of pre-trends, with effects emerging after 2019 and growing thereafter, and are concentrated at the end of the academic year. Youth in exposed areas spent more time studying—roughly 2% more at middle-school age and 5% at high-school age—and 5–6% less on household chores, even during the school closures when formal enrollment was administratively frozen. The effects are visible even after the onset of the COVID-19 pandemic and the associated school closures. Administrative grade-progression records confirm that the contemporaneous reduction in dropout translates into higher educational attainment. The cohort in the second-to-last

grade of high school when the ban hit—which crossed into the final grade in September 2019, after the ban but before the March 2020 school closures—was 6.2 percentage points more likely to reach the final grade in exposed zones (an 11% increase), an *attainment* effect measured free of pandemic confounding. The other cohorts, which crossed into the final grade after the closures, also show a positive, though smaller effect.

More tellingly, the education response is skill-specific. Enrollment in STEM and health programs with chemistry content rose by 5.9% in exposed areas, while non-chemistry STEM fields showed no significant change and overall tertiary enrollment was unchanged. This compositional pattern—a shift toward chemistry-intensive fields with no change in the total—is the distinctive signature of a *skill-specific* shock: a general rise in the returns to schooling would lift enrollment across fields, whereas a chemistry-specific shock reallocates students toward chemistry. It is the composition of education, rather than its level, that most cleanly identifies the mechanism. The effect grows over the event-study horizon as cohorts exposed during secondary school reach university age. Decomposing by gender, women reallocate *within* STEM—from non-chemistry to chemistry-intensive programs—while men’s increased chemistry enrollment appears to draw from non-technical fields or the margin of not attending university, though these gender-specific estimates are less precise.

I show that the demand for chemistry skills rose across the exposed labor markets, suggesting spillovers of the shock to the legal sector through both poaching and input linkages. Exposure to synthetic drug production raised both the price and the quantity of chemistry skill in the legal labor market. The wage premium on a chemistry-relevant degree rose by 6% and employment in chemistry occupations by about 19%. These effects are specific to chemistry: the change in the return to non-chemistry STEM degrees is half as large but statistically insignificant, and the overall wage and employment effects are flat. Price and quantity rising together identify an outward shift in the demand for chemistry skill. Alternative explanations cannot account for this pattern: a supply inflow of chemists would raise employment while lowering the wage; poaching of scarce chemists would raise the wage while cutting employment; a general violence- or displacement-based account would move general quantities without raising the price of the specific skill. Because the legal and illegal sectors hire from the same local pool of chemists, this legal demand is the channel through which the illegal shock is transmitted to reach the schooling choices of all youth rather than just crime-bound ones.

I find suggestive evidence that some skilled chemists were themselves drawn into illegal production. Among workers who completed schooling before the shock, the observed stock of chemistry degree holders in the labor force survey fell by 4.6% in exposed zones, with placebo fields flat, even as observed employment in chemistry occupations rose. Two explanations could account for these missing chemists: they migrated out of exposed zones because of fear of recruitment rather than because of violence or poor economic opportunities, or they remained in the area, moved into the illegal sector, and declined to answer the survey. On the first explanation, I find no evidence of general out-migration. Homicides did not rise with exposure so the violence environment does not

seem to have changed with the shock; the legal return to chemistry rose *within* these zones, so the monetary incentives favored staying; and the out-of-state share of the workforce *fell* slightly rather than rising, whereas an exodus of workers would have pushed it up. Any out-migration, if any, is likely concentrated among the credentialed chemists who were prone to immediate recruitment by the criminal organizations. The other explanation is deliberate concealment by chemists now involved in illegal production, supported by additional evidence from the labor survey. Household refusals to answer the interview rose sharply at the onset of the fentanyl ban: exposure to production raised interview refusals by 1.2 percentage points. There is no effect on non-interviews from vacated dwellings, corroborating that general out-migration did not rise. Fear of victimization is an unlikely driver of these refusals: the shock did not change the violence environment in the short term, and a shock that expands criminal organizations' production without disrupting them is unlikely to cause them to turn to extractive activities like extortion. Also, I do not observe changes in item non-response for income or in income heaping among respondents, which would more clearly signal fear of being targeted for extortion or other extractive crimes.

I observe neither the illegal chemistry premium nor the size of the illegal demand shock. Because the model is invertible, however, I can ask what magnitude of illegal activity would be required to produce the observed rise in the legal chemistry premium—and find that it need not be large. Reproducing that rise requires the illegal chemistry premium to have risen by about a quarter—between 14% and 61% across plausible parameters—and illegal demand for chemistry skill to have grown several-fold to roughly thirtyfold over 2019–2024; a sector employing only a few percent of local chemists produces effects of this size, because its short-run supply is nearly fixed. Independent evidence confirms the illegal sector can reach the necessary size to explain these findings. Most directly, an exposed zone employs on average 675 legal chemistry-related workers; to reach the magnitude that explains the results, the illegal chemists need represent only 1 to 5 percent of them—between 7 and 34 per producing zone, which a handful of clandestine laboratories with a few skilled cooks each would easily provide. The demand growth it requires is easily met: fentanyl seizures rose from almost nothing to more than 1,500 kilograms, ample production growth to account for the demand shift the model gives. Two coarser measures agree on the sector's scale: Mexican cartels are estimated to employ 160,000 to 185,000 people (Prieto-Curiel et al., 2023), disproportionately in the production zones studied here, and the 1.2-percentage-point rise in survey refusals suggests roughly one percent of households were drawn into concealed production.

I subject the findings to extensive robustness checks, the most important concerning how exposure is measured. A laboratory seizure records production and its *detection* at once: where enforcement is weak, production goes unrecorded, so a seizure-based treatment risks measuring the geography of policing rather than of production, carrying potential omitted-variable bias. I address this in two ways. First, I measure where synthesis is *viable* rather than where it is caught, training Random Forest, XGBoost, and ensemble models on pre-2015 fundamentals—proximity to the U.S. border and to the Pacific ports through which precursor chemicals arrive, road connectivity, the local

presence of chemical firms and university chemistry programs, terrain, and cartel presence—while excluding every enforcement- and state-capacity variable. Spatial cross-validation predicts each region from geographically distant ones, so the models cannot learn local policing patterns, and the estimates hold under the resulting predicted and augmented exposure measures. Second, I validate exposure measures against a chemical trace of production that does not depend on policing to be observed: methamphetamine synthesis by the P2P route uses mercury as a catalyst and toluene as a solvent, both of which discharge into local waterways. I find that in national water-quality monitoring the exposed zones carry elevated concentrations of exactly these two compounds.

I also rule out contemporaneous shocks. Treated areas were slightly *less* exposed to Mexico’s 2019 minimum-wage increases, and within control areas minimum-wage exposure *reduced* dropout—the opposite of what would be needed to explain my results. COVID-19 school closures were mandated uniformly at the national level, and effects on study time emerge before the closures. Finally, the 2019 rollback of the conditional cash-transfer program Prospera raised dropout in high-penetration areas—consistent with evidence that such transfers themselves reduce dropout and crime (Chioda et al., 2016)—yet the treatment effect survives controlling for it; strikingly, synthetic-drug exposure offsets *more* of the dropout precisely where Prospera penetration was highest, suggesting that skill-intensive criminal opportunities partially counteract the loss of transfer income.

These findings contribute to several literatures. Most directly, they speak to the economics of crime and the education-crime nexus. The canonical framework posits that education reduces crime by improving legal labor market opportunities (Becker, 1968; Ehrlich, 1973). Empirical work has consistently supported this view: Lochner and Moretti (2004) use compulsory schooling laws to show that education reduces incarceration, Machin et al. (2011) document similar effects in the United Kingdom, and Bell et al. (2022) find that education reduces both violent and property crime. A related literature examines the other direction of causality, how criminal opportunities affect educational investment. Sviatschi (2022) shows that exposure to drug trafficking in Peru increases school dropout and criminal employment among youth. A related strand shows that economic conditions shape criminal participation through the opportunity cost of crime, with falling wages and adverse labor-market shocks raising recruitment and crime (Dube and Vargas, 2013; Dix-Carneiro et al., 2018). Direct evidence that education is rewarded *within* criminal organizations comes from the Italian-American mafia, where the returns to schooling are substantial for members in skill-intensive activities but negligible for those doing physical work (Campaniello et al., 2016). I contribute to this literature by demonstrating that the relationship between criminal opportunities and education depends critically on the *type* of skills that criminal activities demand. When criminal enterprises require skilled labor, the standard result may reverse: criminal opportunities can *increase* educational investment.

Second, the paper connects to the literature on skill-biased technical change and the demand for human capital. A large body of work documents how technological change has increased the relative demand for skilled workers in developed economies (Katz and Murphy, 1992; Autor et al., 2003;

[Acemoglu and Autor, 2011](#)). More recent work examines how skill-biased change affects educational decisions. [Delavande and Zafar \(2019\)](#) show that expected earnings shape university choice, while [Charles et al. \(2018\)](#) document how the housing boom reduced college attainment by raising the returns to low-skill construction work. I extend this literature by showing that skill-biased technical change operates similarly in illegal markets: the shift toward synthetic drug production increased the returns to chemistry skills, and youth responded by investing in relevant education. This represents, to my knowledge, the first documentation of skill-biased technical change in illegal labor markets and its effects on human capital accumulation.

Third, the paper is part of a growing literature on how local labor market shocks affect educational investment. [Atkin \(2016\)](#) shows that export manufacturing opportunities in Mexico reduced educational attainment as youth entered the workforce earlier to take advantage of new jobs. [Adukia et al. \(2020\)](#) demonstrates that road construction in India increased school enrollment by improving access to labor markets where education is rewarded. [Shah and Steinberg \(2017\)](#) shows that rising agricultural wages in India reduced schooling as the opportunity cost of education increased. I contribute to this literature by examining how *illegal* labor market shocks reshape educational decisions, showing that the mechanisms documented in legal labor markets—changes in opportunity costs and returns to specific skills—operate similarly when the labor demand originates from criminal enterprises.

Fourth, this research adds to the list of applications of machine learning methods in economics. A growing literature uses prediction-focused methods to address measurement challenges in causal inference ([Mullainathan and Spiess, 2017](#); [Athey and Imbens, 2019](#)). [Kleinberg et al. \(2015\)](#) introduce the concept of “prediction policy problems” where the goal is to predict an outcome that is only partially observed. [Kleinberg et al. \(2018\)](#) apply this framework to bail decisions, where observed outcomes depend on prior judicial choices. In the context of illegal activities, [de Roux et al. \(2018\)](#) and [Battaglini et al. \(2025\)](#) use machine learning to predict tax evasion and improve audit targeting; closest to my setting, [Campedelli et al. \(2024\)](#)—in parallel work—predict mafia infiltration of Italian municipalities, observed only through state-ordered council dismissals, and use the predicted index to flag undetected cases. I contribute to this literature by using machine learning to address a fundamental measurement problem in studying illegal markets: observed enforcement actions (laboratory seizures) reflect both underlying illegal activity and enforcement effort. By training models on features that predict production capacity while excluding enforcement-related variables, and using spatial cross-validation to prevent learning enforcement patterns, I construct treatment measures that better capture latent production. Because the label itself reflects enforcement, I do not optimize a single classification threshold—which would risk fitting where the state happened to seize laboratories rather than where production occurs—but instead report estimates across a range of thresholds and treatment definitions, treating the resulting spread as a robustness band.

Finally, this work contributes to understanding the broader economic effects of organized crime, a literature spanning the economics of illegal drug markets and enforcement ([Mejía and Restrepo,](#)

2016) and the governance of territory by criminal organizations (Blattman et al., 2025). Beyond the well-documented effects on violence (Dell, 2015; Castillo and Kronick, 2020), migration (Daniele et al., 2023), and political outcomes (Alesina et al., 2019), criminal organizations reshape local economies through labor market channels. Pinotti (2015) estimates the economic costs of organized crime in Italy, while Mirenda et al. (2022) document how mafia presence distorts firm behavior. I show that criminal influence extends to human capital investment decisions among all youth in affected communities—not only those who ultimately pursue criminal careers—through the educational incentives that illegal labor markets create and their spillovers to the legal market. This suggests that the social costs of organized crime include distortions to human capital accumulation that may persist long after criminal organizations themselves are dismantled.

The remainder of this paper proceeds as follows. Section 2 provides background on drug trafficking in Mexico, the drivers of the transition to synthetic drug production, and its labor market implications. Section 3 develops the theoretical framework—a local market for chemistry skill—that generates the paper’s testable predictions and that can be inverted to gauge the size of the illegal demand shift. Section 4 describes the data sources and construction of key variables, and Section 5 presents the empirical strategy. Section 6 reports the main results on dropout, grade progression, and tertiary enrollment. Section 7 turns to the labor-market mechanism: it establishes the low-violence, low-enforcement environment and the concealment of illegal workers, then documents the joint rise in the price and quantity of chemistry skill, the de-skilling of the observed workforce, and the disappearance of credentialed chemists from observation. Section 8 uses the model’s invertibility to gauge the size of the illicit sector and benchmark it against external evidence. Section 9 presents robustness checks addressing alternative treatment definitions, sample composition, and contemporaneous shocks, and Section 10 concludes.

2 Context

2.1 Mexico’s Drug Trafficking Industry

Mexico is the primary source and transit country of illegal drugs entering the United States, the world’s largest consumer market. The scale of this industry is enormous: the United Nations Office on Drugs and Crime estimates that inward illicit financial flows from trafficking heroin, cocaine, and methamphetamine averaged roughly \$12 billion per year over 2015–2018—comparable to the value of Mexico’s agricultural exports over the same period—while U.S. and Mexican government estimates place the drug-trafficking proceeds flowing from the United States into Mexico at \$19–29 billion annually (United Nations Office on Drugs and Crime, 2023; U.S. Immigration and Customs Enforcement, 2010). Major drug trafficking organizations, often called “cartels,” control production, transportation, and distribution networks across Mexico and involve sophisticated logistics, financial operations, and enforcement capabilities.

Historically, Mexican drug production centered on plant-based drugs like marijuana and heroin.

Cannabis and opium poppy cultivation has occurred in Mexico for decades. Both crops are grown primarily in mountainous regions with suitable climate and terrain, particularly in states like Sinaloa, Durango, and Guerrero. Production required agricultural knowledge, access to land, and willingness to work in remote and often dangerous locations.

2.2 The First Synthetic Drug Expansion: methamphetamine

Beginning in the late 2000s and accelerating through the 2010s, Mexican trafficking organizations increasingly diversified towards synthetic drug production, particularly methamphetamine (Figure 1). This transition was driven by several factors: disruptions to the availability of precursor chemical supplies in the United States that had previously supported domestic meth production; growing demand from American consumers; and the higher profit margins available from synthetic drugs, which require no land and can be produced year-round wherever the necessary inputs can be procured (Asmann, 2021).

Methamphetamine production requires fundamentally different capabilities than agricultural drug production. Synthesizing methamphetamine involves chemical reactions, precursor procurement, laboratory setup and operation, and quality control. While “shake and bake” methods allow small-scale production with minimal expertise,¹ the large-scale operations that supply wholesale markets require genuine chemical knowledge: understanding reaction kinetics, handling volatile and toxic substances, troubleshooting production problems, and adapting methods when precursor availability changes.

Mexican seizure data reflects this expansion. Figure 1a shows that between 2005 and 2015, authorities dismantled an increasing number of clandestine laboratories, predominantly producing methamphetamine. The geographic concentration of these seizures, shown in Figure 1c, reveals that synthetic drug production was not uniformly distributed: only certain regions developed specialized production capabilities, particularly in areas with proximity to precursor chemical supply routes through Pacific coast ports, U.S. border access, and historical presence of criminal organizations.

2.3 The Second Synthetic Drug Expansion: fentanyl

The opioid crisis in the United States created the demand conditions that would result in the drug production transformation studied in this paper. Beginning in the 1990s, aggressive marketing of prescription opioids led to widespread addiction (Van Zee, 2009). When the 2010 reformulation of

¹The “shake and bake” or “one-pot” method combines pseudoephedrine with volatile reagents—typically lithium stripped from batteries, a source of ammonia such as ammonium nitrate, and a flammable solvent—inside a single sealed container that is shaken to drive the reaction. It requires little equipment and only small quantities of precursor, but the reaction is prone to fire and explosion, has a high failure rate, and yields small batches of low-purity product. It is therefore ill-suited to the sustained, high-volume output required to supply wholesale and export markets, which is why the industrial-scale production discussed here relies on trained personnel rather than amateur “cooks” (National Drug Intelligence Center, 2009).

OxyContin and related regulatory crackdowns restricted access to prescription pills, users turned to illegal substitutes—first heroin (Alpert et al., 2018; Evans et al., 2019), then, around 2015, fentanyl.

Fentanyl—a synthetic opioid 50-100 times more potent than morphine—offered enormous advantages to suppliers: no agricultural land required, year-round production, and extreme potency meaning small shipments yielded large quantities of street-level product (National Institute on Drug Abuse, 2021). Chinese legal chemical manufacturers, operating with minimal regulatory oversight, produced fentanyl and intermediaries shipped it directly to American consumers through the postal system, often in packages small enough to evade detection (U.S. Senate Permanent Subcommittee on Investigations, 2018). By 2016, fentanyl had become the leading cause of overdose deaths in the United States (Figure 2).

On April 1, 2019, responding to sustained diplomatic pressure from the United States, China announced a ban on fentanyl and its analogues (the *fentanyl ban*), effective May 1, 2019 (Congressional Research Service, 2019).² The ban disrupted supply but did not eliminate American demand. Mexican trafficking organizations, already operating methamphetamine laboratories and controlling distribution routes into the United States, seized the opportunity: China’s role shifted from exporting finished fentanyl to supplying precursor chemicals to Mexico, which Mexican organizations now synthesize domestically (U.S. Drug Enforcement Administration, 2020) and then smuggle through the U.S.-Mexico border. This fentanyl expansion built on existing methamphetamine infrastructure—trained personnel, supplier relationships, laboratory equipment, and distribution channels—giving established meth-producing regions a decisive advantage in expanding to fentanyl.

Figure 1 illustrates these trends. Panel A shows clandestine laboratory seizures by the Mexican government, which fluctuated between roughly 100 and 250 per year through the 2010s before rising sharply after 2020, to nearly 950 by 2023. Panel B shows fentanyl seizure volumes, which climbed from negligible levels before 2017 to about 340 kilogram-equivalents in 2018 and, after a temporary dip in 2019, grew rapidly to over 1,500 by 2023. Notably, fentanyl did not replace methamphetamine: the two are pharmacologically distinct—one a synthetic opioid, the other a stimulant—but in the market they have proven complementary rather than substitutable, with methamphetamine production continuing to grow alongside fentanyl and the two increasingly consumed together.³

²Formally, China placed all fentanyl-related substances under national control as a single class (“class scheduling”), rather than banning individually named compounds one by one. This brought every current and future variant of the fentanyl molecule under control at once, closing the loophole by which producers had evaded regulation by making minor chemical modifications to synthesize novel, technically legal analogues.

³The share of U.S. overdose deaths involving both fentanyl and a stimulant rose from under 1% in 2010 to roughly one-third by 2021, making stimulants the most common drug class found alongside fentanyl in fatal overdoses (Friedman and Shover, 2023).

2.4 Labor Market Implications

These organizations are not only financially large but also major employers of labor. Combining data on homicides, disappearances, and incarcerations, [Prieto-Curiel et al. \(2023\)](#) estimate that Mexican criminal organizations collectively employ between 160,000 and 185,000 people—placing them, in aggregate, among the country’s five largest employers—and must recruit at least 350 new members every week simply to offset losses from arrests and inter-cartel conflict. The scale of this workforce, and the constant recruitment needed to sustain it, means that shifts in what these organizations produce translate directly into shifts in the kind of labor they demand.

The shift toward synthetic drug production creates demand for chemistry skills—individuals with knowledge of chemical reactions, laboratory procedures, and scientific reasoning—that has been documented by various sources. Documents from Mexico’s Secretariat of National Defense, leaked in 2022, revealed that the Sinaloa Cartel was actively recruiting chemists.⁴ Chemistry professors in Baja California have reported fearing recruitment attempts, with some avoiding mention of their profession to prevent attracting criminal attention ([Dudley, 2023](#)).⁵

One might ask whether trafficking organizations truly need many trained chemists, or whether a few experts could simply train others on the job. The evidence suggests that skill requirements are substantial and difficult to substitute. Fentanyl synthesis is demanding at every stage—some routes require specific atmospheric conditions, and converting pre-precursors into precursors can be fatal if performed incorrectly ([Dudley, 2023](#))—and the further back along the route producers must work, the more knowledge, equipment and precision are required. Producers are increasingly pushed into more chemically demanding routes as precursor availability is disrupted.⁶ What organizations require is therefore not a one-time transfer of a recipe but a standing capacity to re-engineer production. Untrained “cooks” are costly on both margins: accidents destroy product and personnel, while inconsistent dosing in counterfeit pills attracts law enforcement attention and damages market reputation.⁷

This skill-biased shift in illegal markets carries at least three implications. I state each as a potential mechanism, formalized in the model in Section 3, and pair it with the qualitative evidence—

⁴The documents were part of the “Guacamaya Leaks” hack of SEDENA servers. See [Garcia \(2023\)](#).

⁵One professor recounted being approached in the 1990s to help produce methamphetamine. In 2021, a chemistry student in Tijuana was killed months before graduation. Prosecutors investigated whether the attack was retaliation for refusing to work for criminal groups. ([Garcia, 2023](#))

⁶NPP and ANPP, the immediate fentanyl precursors, were added to Table I of the 1988 Convention in 2017; producers shifted routes, prompting the scheduling of 4-AP, 1-boc-4-AP and norfentanyl in 2022 and of 4-piperidone and 1-boc-4-piperidone in 2024. Ten substances were added to Table I in the eight years to 2022, several designer precursors with no known legitimate use, which the INCB attributes to how quickly traffickers substitute uncontrolled chemicals for controlled ones ([United Nations Office on Drugs and Crime, 2022](#); [International Narcotics Control Board, 2023](#)).

⁷Early U.S. fentanyl-pill operations blended fentanyl powder with inert fillers in household blenders, distributing the drug unevenly—the “chocolate chip cookie” effect, whereby some pills contain little fentanyl and others a lethal dose. Reporting links this crude mixing, as much as deliberate potency, to many fatal overdoses ([Quinones, 2021](#)); DEA testing finds roughly six in ten seized counterfeit pills contain a potentially lethal dose ([U.S. Drug Enforcement Administration, 2022](#)).

reporting, interviews, and government documents—consistent with it. First, the shift potentially raises the returns to chemistry-related skills for individuals willing to work in illegal production. Consistent with this, a chemistry professor in Sinaloa told InSight Crime that he has had multiple students eager to enter synthetic drug production, some enrolling in his course specifically for that purpose; as their numbers grew, he stopped trying to dissuade them and began teaching laboratory-safety techniques to reduce accidents (Dudley, 2023).

Second, this demand may spill over to legal chemistry-related sectors through two channels: trafficking organizations compete for chemically skilled workers with legitimate employers—a *poaching* channel—and illicit production draws on *input linkages* to legal firms that supply precursors, equipment, and expertise. Consistent with the input-linkage channel, criminal organizations establish front companies in the chemical and pharmaceutical sectors to launder money and procure precursors.⁸ My field interviews suggest these linkages can run deeper still: informants described operators who own a clandestine laboratory alongside a registered legal firm that both sources hard-to-obtain inputs for the lab and enrolls the lab’s workers in the social security system, so that employees of an illegal operation accrue health insurance and pension contributions.⁹ Such arrangements blur the boundary between licit and illicit activity: a single firm serves both markets, and the infrastructure and expertise built for illegal production can seed genuinely legal enterprises.

Third, and most consequential for my analysis, these spillovers imply that a criminal demand shock can reshape educational investment well beyond those who enter the trade. Because legal and illegal employers bid for chemistry skill from the same local pool, a rise in illegal demand raises the return to chemistry education for all youth—not only those bound for the illegal sector—who may therefore invest in chemistry skills even if they never intend to work illegally.

3 Theoretical framework

To understand how a shock to the illegal demand for chemistry skill could shape educational choices, I use a Roy-type sorting model of a local labor market. Youth sort across two sectors, a *legal* one and an *illegal*, drug-producing one, and choose whether and in what field to be educated; the return to chemistry skill is set by clearing a local market for chemistry-skilled labor in which both sectors bid. The model does two things. It generates the testable predictions that the rest of the paper examines. Also, it disciplines the natural objection that clandestine production is too small a part of the local economy to move wages or schooling, by showing this depends on *elasticities* rather than headcounts. Using the observed wage and educational responses to the shock, the model can then be inverted to gauge the size of the illegal demand shift. I keep only the structure needed for these purposes here; the full model—decision rules, market clearing, existence and uniqueness,

⁸In October 2025, the U.S. Treasury sanctioned twelve Mexico-based companies affiliated with the Sinaloa Cartel’s Los Chapitos faction, accusing them of using pharmaceutical, laboratory, chemical, cleaning, and real estate businesses to supply precursor chemicals for fentanyl production (U.S. Department of the Treasury, 2025).

⁹Author field interview, Sinaloa, June 2023.

and comparative statics—is developed in Appendix B, with its elements summarized in Appendix Table B.1.

3.1 Setting

Consider the youth of a local labor market, empirically a commuting zone. Each makes three choices: whether to complete an additional level of schooling, $e \in \{0, 1\}$, at cost c ; if so, a field of study $f \in \{C, N\}$ —chemistry-intensive (C) or not (N)—at an idiosyncratic net cost η_i for C ; and a sector of employment $s \in \{L, I\}$, legal (L) or illegal (I). Illegal work pays more at the base ($w_{I0} > w_{L0}$) but carries a risk of arrest $p(\theta)$ that is decreasing in an individual’s type θ : a higher θ means a lower chance of being caught—for instance, family or social ties to a trafficking organization that bring trust and protection, or the local know-how to operate without drawing law-enforcement attention.¹⁰ Earnings are

$$w_s(e, f) = w_{s0} + e \cdot [\pi_s \mathbf{1}\{f = C\} + \beta_s^N \mathbf{1}\{f = N\}], \quad s \in \{L, I\}. \quad (1)$$

The non-chemistry returns β_s^N are parameters, but the two *chemistry premiums* π_L and π_I are prices, determined by clearing a local market for chemistry skill. Allowing them to differ is essential: because illegal work carries arrest risk, the illegal sector must pay a compensating differential to attract chemists, so in equilibrium $\pi_I > \pi_L$, and the size of that wedge is itself an equilibrium object.

Choices are cutoff rules in (θ, η_i) . Given schooling and field, the sector choice compares the legal wage with the risk-adjusted illegal payoff,

$$\text{work illegally} \iff (1 - p(\theta)) w_I(e, f) - p(\theta) P > w_L(e, f), \quad (2)$$

where P is the penalty on arrest, so lower-risk (high- θ) individuals sort into illegal work; and the chemistry field is chosen when its premium, net of the extra cost η_i , exceeds the non-chemistry return. Higher premiums push both margins—more schooling, and more of it in chemistry. The thresholds are solved explicitly in Appendix B; here they matter only through the equilibrium premiums.

The legal and illegal sectors are connected by labor poaching and input linkages. They hire chemists from the *same local pool*: a chemist indifferent between them equates risk-adjusted pay at the margin, so the two sectors are gross substitutes in chemistry skill—an illegal boom that bids up π_I pulls marginal chemists across the sorting threshold and, by making them scarce, raises π_L . Illegal production also demands other chemistry-related *inputs*—precursors, equipment,

¹⁰The arrest-risk interpretation is not essential. The type θ could equally represent an individual’s disutility or distaste for working in the illegal sector, or any other one-dimensional trait that raises the relative payoff of illegal work; all the model requires is that individuals differ along a single dimension that orders them by their propensity to work illegally. We adopt the arrest-risk formulation because it lets the penalty P and enforcement enter the payoff transparently.

and laboratory and commercial services—so an illegal boom can also shift the *legal* demand for chemistry skill outward. Writing sector demands in iso-elastic form,

$$H_L^d = B \pi_L^{-\sigma_L}, \quad H_I^d = A \pi_I^{-\sigma_I}, \quad (3)$$

I summarize the input link by a pass-through $\phi \equiv d \ln B / d \ln A \geq 0$, where $\phi = 0$ is a pure poaching scenario. China’s 2019 fentanyl ban, by making domestic synthesis from uncontrolled precursors the profitable production technology, is an outward shift in illegal demand—an increase in A in exposed commuting zones. Market clearing in (π_L, π_I) has a unique solution (Appendix B); the shock’s effects follow by comparative statics.

3.2 Predictions

The shock $d \ln A > 0$ raises the equilibrium premiums and, through the cutoff rules, the education outcomes. Four predictions follow, each with a direct empirical counterpart.

1. **Both chemistry premiums rise, the legal one by less.** The illegal premium rises directly; the legal premium rises too, and—through the poaching link alone—by a strictly smaller amount, $0 < d\pi_L/d\pi_I < 1$. *Empirical counterpart:* the rise in the legal-market wage premium for chemistry-course degrees (Table 6, Panel A).
2. **Legal chemistry *employment* rises only if the supply link is active.** Under pure poaching ($\phi = 0$), legal chemistry *wages* rise while legal chemistry *employment falls*, because the illegal sector draws chemists away along the legal demand curve; legal employment rises if and only if $\phi > 0$. The joint finding that exposure raises both the legal premium *and* legal chemistry employment therefore identifies $\phi > 0$: the boom shifted legal demand outward, it did not merely poach. *Empirical counterpart:* the rise in chemistry-occupation employment (Table 6, Panel B).
3. **Educational investment rises.** A higher π_I lowers the schooling threshold for illegal-bound (high- θ) youth—a direct effect—and the induced rise in π_L lowers it for everyone—a spillover—so dropout falls and enrollment rises. *Empirical counterpart:* the dropout and enrollment results of Section 6.
4. **Field choice shifts toward chemistry.** Because only the *chemistry* premiums rise, the field cutoff moves toward C : chemistry-intensive enrollment rises while non-chemistry and total enrollment do not. *Empirical counterpart:* the compositional enrollment result (Table 4).

3.3 Sizing the Illegal Demand Shift

Because both sectors clear a single local market for chemistry skill, the legal wage response the design estimates is quantitatively informative about the illegal shock itself. Over the study horizon

the education pipeline cannot add chemists—producing one takes years of schooling—so the local stock is effectively fixed, and every chemist the illegal sector hires is one the legal sector loses. Moving along the legal demand curve (3), the legal premium then satisfies

$$d \ln \pi_L = \frac{1}{\sigma_L} \frac{dH_I}{H_L}, \quad (4)$$

so the estimated legal premium response equals illegal hiring *as a share of local legal chemistry employment*, scaled by the inverse legal demand elasticity. The effect is therefore governed by elasticities rather than by the sector’s raw size: with specialized legal demand (σ_L small), an illegal sector that hires only a few percent of local chemists can move the legal premium by several percent. The same structure makes the model *invertible*: the estimated legal premium, together with defensible elasticities, recovers the two objects the data cannot show directly—the change in the illegal chemistry premium and the size of the illegal demand shift. I carry out that inversion, and confront it with external evidence on laboratory prevalence, in Section 8. Throughout, this exercise is a consistency check on the mechanism, not a source of identification; identification rests on the empirical strategy described in 5.

4 Data

The analysis combines data from several sources: labor force surveys measuring educational outcomes, administrative records on academic enrollment, military records on drug seizures and laboratory dismantlement, and geographic data for constructing local labor markets. This section describes each data source and the construction of key variables.

4.1 Educational Outcomes

School Dropout (ENOE). The primary measure of educational attainment comes from Mexico’s National Survey of Occupation and Employment (ENOE), a quarterly rotating panel that is the country’s principal labor force survey. ENOE provides sociodemographic information, including current enrollment for all members of the household, as well as labor market information for members 12 years or older.

I measure dropout using a school-age cohort approach based on expected grade progression. Mexico’s school entry rules determine that children born before a specific cutoff date should start primary school at age six. Using birth dates, I identify individuals who should be enrolled at each educational level if they progressed on time: primary school age (6–11), middle school age (12–14), and high school age (15–17). Dropout is defined as not currently enrolled in school among individuals of the relevant school age.¹¹

¹¹This approach captures whether school-age individuals are in school, regardless of their actual grade level. A 13-year-old who is enrolled but repeating primary school is not counted as a dropout, while a 13-year-old not enrolled at all is.

The analysis uses data from 2014 to 2024, spanning five years before and after China’s fentanyl ban in April 2019. I exclude summer quarters (Q3) due to high seasonality in enrollment—many students report not being enrolled during vacation periods, which would confound the dropout measure. Figure A.1 illustrates this pattern. I also exclude the second quarter of 2020, when ENOE temporarily switched to a phone-based survey due to COVID-19 restrictions, resulting in a non-comparable sample.¹² The resulting analysis sample covers 1,263 commuting zones.

Tertiary Enrollment (ANUIES). For tertiary (university) enrollment, I use administrative data from the National Association of Universities and Higher Education Institutions (ANUIES), which collects annual enrollment statistics from all Mexican higher education institutions. These data provide enrollment counts by institution, campus, and academic program, but do not include information on students’ municipality of origin. This creates a challenge: enrollment is observed at the destination (where universities are located), but treatment is defined at the origin (students’ home commuting zones).

To assign enrollment to origin locations, I combine ANUIES data with migration patterns from the 2010 Population Census (10% sample). The census asks respondents where they lived five years prior, allowing me to construct an origin-destination matrix for individuals of tertiary education age (18–24). For each destination CZ d , I compute the share of tertiary-age individuals who originated from each origin CZ o , denoted $P_{d,o}$. I then allocate enrollment at each destination back to origin CZs using these shares:

$$\tilde{Y}_{o,t} = \sum_{d=1}^N P_{d,o} \cdot Y_{d,t} \quad (5)$$

where $Y_{d,t}$ is observed enrollment at destination d in year t , and $\tilde{Y}_{o,t}$ is the imputed enrollment originating from CZ o .

This procedure allows me to measure how tertiary enrollment responds to local conditions in students’ home commuting zones, rather than where universities happen to be located. For commuting zones without a university campus, this is the only way to capture tertiary enrollment responses. For commuting zones with campuses, the measure captures both local students and those who migrate elsewhere for university.

Figure A.2 characterizes the student migration matrix across commuting zones. Panel (a) shows that the average origin CZ sends students to 5.73 destination CZs, with a probability-weighted mean migration distance of 166 kilometers. Panel (b) reveals that migration probabilities are heavily right-skewed: the vast majority of origin-destination pairs have near-zero migration flows (diagonal elements representing students who stay approach 100%, while most off-diagonal migration probabilities cluster near zero). Panel (c) focuses on the off-diagonal migration probabilities excluding near-zero values, showing that conditional on migration occurring, flows are approximately

¹²INEGI resumed in-person interviews in Q3 2020, though with modified protocols. Results are robust to excluding the entire year of 2020.

log-normally distributed with most probabilities in the 0.01–0.20 range.

The key identifying assumption is that migration shares are stable conditional on treatment. If treatment alters migration patterns, the interpretation of the estimates changes. Consider three scenarios. First, if students from treated origins increasingly migrate to treated destinations—perhaps because treated areas develop stronger chemistry programs or labor market networks—the 2010-based weights would under-allocate enrollment to these destinations, biasing estimates toward zero and making the results conservative. Second, if treatment increases migration to control destinations—for instance, if students seek distance from criminal activity—the measure would over-allocate enrollment to treated origins, potentially generating spurious positive effects. Third, if treatment reduces migration overall because local opportunities (legal or illegal) reduce incentives to leave, the fixed weights would over-allocate enrollment to distant destinations, again attenuating the estimates.

The descriptive patterns in Figure A.3 help assess these concerns. The figure shows that cross-treatment migration is limited: students from treated origins overwhelmingly migrate to treated destinations (91.3%), while students from control origins predominantly stay within control areas (99.7%), with only 0.3% migrating to treated destinations. This limited cross-treatment flow has important implications for potential bias. First, enrollment changes in one treatment group are unlikely to contaminate the other through reallocation—if treated origins increase enrollment, this will primarily show up at treated destinations and be correctly allocated back to treated origins. Second, if treatment induces students from treated origins to shift toward different treated destinations—for instance, toward CZs with stronger chemistry programs—the 2010 weights would misallocate enrollment across treated CZs, but the overall enrollment attributed to treated origins as a group would remain roughly correct. Third, the most concerning scenario—treatment increasing migration from treated to control destinations, which would spuriously inflate treated-origin enrollment—is unlikely given the small baseline share (8.7%) flowing in this direction. If anything, treatment effects that strengthen local opportunities might reduce out-migration to control areas, which would cause the 2010 weights to over-allocate enrollment to control destinations, attenuating the estimated effects. More broadly, evidence suggests that tertiary migration is primarily driven by university supply and amenities rather than local labor market conditions (Böckerman and Haapanen, 2013), though drug violence has been shown to affect skilled migration in Mexico (Aldeco Leo et al., 2024).

To examine education choices related to the specific skill content of the shock, I classify academic programs by their chemistry content using detailed curriculum analysis. Drawing on publicly available curricula (*mapas curriculares*) from the National Autonomous University of Mexico (UNAM), Mexico’s largest and most prestigious public university, I systematically measure the chemistry intensity of different academic programs. I then map all programs in the ANUIES tertiary enrollment database to these UNAM reference programs to assign chemistry intensity measures to each

program nationwide.¹³

Using this approach, I classify academic programs into three categories based on their chemistry course requirements: STEM+H programs with chemistry, STEM+H programs without chemistry, and non-STEM+H programs. STEM+H fields include Natural Sciences, Mathematics and Statistics; Engineering, Manufacturing and Construction; Agriculture, Forestry, Fisheries and Veterinary; and Health and Welfare, following the International Standard Classification of Education (ISCED) 2013 framework.

Programs classified as STEM+H with chemistry are those that require at least one chemistry course according to their matched UNAM curriculum. These include chemical engineering, pharmaceutical sciences, biotechnology, biochemistry, food science engineering, chemistry, biology, and broad engineering degrees like industrial engineering. Among STEM+H programs, 48.2% require at least one chemistry course, with a mean of 1.96 chemistry courses (ranging from 0 to 27). STEM+H programs without chemistry include science, technology, engineering, and health programs where chemistry is absent from the curriculum—for example, architecture, civil engineering, computer science, mathematics, and mechatronics engineering. Medicine and nursing are classified as STEM+H with chemistry as they require introductory chemistry coursework.¹⁴

A related limitation to correctly estimate the effects on specific degrees is that the migration matrix captures overall flows of tertiary-age students, not field-specific migration. Yet I examine effects on chemistry versus non-chemistry enrollment separately. This requires assuming that migration patterns do not vary systematically by field of study. If treatment increases demand for chemistry skills, students from treated origins may be more willing to migrate further or to different destinations to access chemistry programs—particularly if such programs are concentrated in a few major universities. By applying the pre-treatment overall migration weights, I would under-allocate this new chemistry enrollment to treated origins (assigning it instead based on historical patterns that reflect the average student, not the marginal chemistry student). This implies that my estimates of increased chemistry enrollment from treated origins are likely attenuated, making the findings conservative. Reassuringly, the share of chemistry programs does not vary significantly across treated and control destinations (Figure 4), suggesting that students from treated origins need not systematically alter their migration patterns to access chemistry education.

Grade Progression (Formato 911). The ENOE measures whether a school-age individual is currently attending school, but not whether students advance through and complete each level. To measure grade progression directly, I use *Formato 911*, the administrative census of the Mexican school system. Collected by the Ministry of Public Education (SEP) at the start of every school year, it records enrollment, first-time entrants, and repeaters by grade for the universe of

¹³Appendix D describes the full classification methodology in detail.

¹⁴Figure D.1 in the Appendix shows the full distribution of chemistry intensity across programs. Figure D.2 lists the most common programs in each category.

schools in the country. From these counts I construct grade-to-grade continuation rates for each commuting zone and follow successive entry cohorts through the system, measuring the share of students who enter middle school and remain enrolled through the final grade of high school—a six-grade trajectory. Because the measure is built from grade-specific enrollment aggregated to the commuting zone, transfers between schools within a zone net out. The data span the 2010–2011 through 2023–2024 school years, cover the same commuting zones as the ENOE sample, and are assembled from two administrative sources over the sample period; I subject their concatenation to extensive harmonization checks. This administrative measure complements the survey: whereas the ENOE captures contemporaneous enrollment at a point in time, Formato 911 captures whether a cohort actually advances to the terminal grade of high school—an attainment margin that a contemporaneous enrollment snapshot need not reflect.

4.2 Treatment Variables

The empirical strategy compares commuting zones with high versus low exposure to the expansion of synthetic drug production following China’s 2019 fentanyl ban. High-exposure areas are those most likely to expand production in response to the ban because they possess the infrastructure, expertise, and supply networks required for synthetic drug manufacturing.

I define exposure based on pre-existing methamphetamine production capacity. As described in Section 2, Mexican trafficking organizations developed domestic methamphetamine production during the first synthetic drug expansion in the 2000s and 2010s. Commuting zones that established meth laboratory operations during this period had accumulated the key inputs for fentanyl production: clandestine laboratories, trained chemists and cooks, precursor chemical supply chains, and relationships with trafficking organizations. When China’s ban disrupted the fentanyl supply chain in 2019, these areas had a comparative advantage in pivoting to domestic fentanyl synthesis.

Clandestine Laboratory Seizures (SEDENA). Data on synthetic drug production infrastructure comes from Mexico’s Secretariat of National Defense (SEDENA), which maintains records of military operations including the dismantlement of clandestine drug laboratories. These records cover the period 2000–2023 and include the number of laboratories in each municipality per year.

The main treatment variable is a binary indicator equal to one if any methamphetamine laboratory was seized in the commuting zone before 2015. This cutoff ensures that treatment is determined by infrastructure from the first synthetic drug expansion rather than the post-2019 fentanyl expansion.

Figure 3 maps the spatial distribution of treated and control commuting zones, along with subsequent drug production activity. The map reveals strong geographic persistence in synthetic drug production: post-2019 clandestine laboratory seizures (shown as green diamonds) occur predominantly in areas that already had pre-2015 exposure, with a correlation of 0.42 between the exposure variable and post-2019 lab presence. Fentanyl seizures (purple circles) show a somewhat different spatial pattern, with a correlation of 0.21 with pre-2015 exposure. This partial divergence reflects

that fentanyl seizures typically occur later in the distribution chain and closer to the U.S. border, capturing final product seizures rather than production sites. The concentration of both pre- and post- production in western and central Mexico reflects proximity to Pacific ports where precursor chemicals arrive and to northern border regions where synthetic drugs are ultimately trafficked into the United States.

4.3 Machine Learning Prediction of Latent Production

A concern with using seizure-based treatment measures is that observed seizures capture the intersection of two processes: actual production and enforcement effort. Areas with intensive law enforcement may have more seizures despite similar underlying production; areas with weak enforcement may have production that goes undetected.

To address this, I develop a machine learning approach to predict latent production potential independent of detection. This approach draws on a growing literature applying machine learning to policy problems in economics. [Athey and Imbens \(2019\)](#) provide an overview of machine learning methods for economists, while [Kleinberg et al. \(2015\)](#) introduce the concept of “prediction policy problems”—settings where the key challenge is prediction rather than causal inference, and where machine learning methods can substantially outperform traditional approaches. My application fits this framework: I seek to predict latent production capacity from observable characteristics, which is fundamentally a prediction task. Similar to [Kleinberg et al. \(2018\)](#), who use machine learning to predict defendant behavior in bail decisions where observed outcomes depend on prior judge decisions, my setting involves predicting concealed activity (drug production) where observed outcomes (seizures) reflect both underlying activity and enforcement effort. Related work applies machine learning to detect tax fraud, where observed audits similarly capture both evasion and enforcement intensity ([Battaglini et al., 2025](#)).

I train Random Forest, XGBoost, and ensemble models on features that predict where production is viable—not only where authorities detect it. These include economic fundamentals (distance to U.S. border, road network connectivity, proximity to Pacific ports, presence of chemical firms and university chemistry programs), demographics (population structure, schooling levels), geography and climate (bioclimatic variables, elevation, terrain ruggedness), and indicators of illegal activity (cartel presence, other drug seizures, agricultural suitability for drug crops). All features are measured in 2015 or earlier. Importantly, I exclude variables related to enforcement intensity or government capacity, which would predict detection rather than production.

To ensure the model learns production patterns rather than enforcement patterns, I implement spatial cross-validation: Mexico is partitioned into five geographic folds, and in each iteration the model trains on four folds and predicts on the held-out fold. This forces predictions for each region to be based on patterns from geographically distant areas with potentially different enforcement regimes.

I construct two types of alternative treatment measures. *ML-predicted* treatment classifies CZs based solely on predicted probabilities at the 25th or 50th percentile threshold. *Augmented* treatment combines baseline seizures with ML predictions: $\hat{D}_i^{Aug} = \max\{D_i, \hat{D}_i^{ML}\}$, ensuring all observed seizures remain treated while adding likely producers that evaded detection. The baseline identifies 163 treated CZs; augmented measures at the 25% threshold identify 270–277 CZs, adding 107–114 CZs with high predicted production potential but no observed seizures. Appendix C provides full details on the predictor data, model specification, feature importance, and performance metrics.

4.4 Commuting Zones

I aggregate the analysis to commuting zones (CZs) rather than municipalities for several reasons. First, labor markets typically span municipal boundaries, particularly in metropolitan areas. Second, aggregation reduces measurement error in treatment status: a single municipality missing a lab seizure may simply have been lucky to avoid detection, but a commuting zone without any seizures more plausibly lacks production infrastructure. Third, commuting zones provide appropriate geographic units for analyzing educational choices, as students often attend schools outside their municipality of residence.

I construct CZs for Mexico following the methodology of Duranton (2015), originally developed to delineate metropolitan areas in Colombia. Using the 2010 Population Census, which includes information on both municipality of residence and municipality of work, I construct origin-destination commuting matrices. Mexico’s 2,478 municipalities are then aggregated into CZs based on commuting flows, resulting in 1,776 commuting zones, of which 1,226 are covered in the analysis sample.¹⁵

4.5 Water Quality Data

To validate the clandestine drug production exposure measures, I use water quality monitoring data from Mexico’s National Water Quality Monitoring Network (RENAMECA). RENAMECA collects samples at approximately 5,000 monitoring sites across Mexico’s major water bodies, measuring concentrations of heavy metals, organic solvents, and other pollutants. I focus on chemicals associated with methamphetamine synthesis using the P2P method: mercury (used as a catalyst in the aluminum-mercury amalgam reduction step) and toluene (used as a solvent for extraction and purification). I aggregate site-level measurements to commuting zones using inverse-distance weights from each monitoring site to the commuting zone centroid, constructing annual pollution indices for the period 2013–2020. Appendix F describes the chemical signatures of drug production and presents the full validation analysis.

¹⁵The remaining commuting zones were not sampled by ENOE during the analysis period.

5 Empirical strategy

The empirical strategy compares educational outcomes in commuting zones with pre-existing synthetic drug production infrastructure (treated) to those without (control), before and after China’s 2019 fentanyl ban. I estimate separate specifications for dropout (individual-level) and tertiary enrollment (commuting zone-level).

5.1 Dropout before Tertiary Education

For dropout outcomes, I estimate the following difference-in-differences specification at the individual level:

$$y_{i,c,t} = \beta \cdot Exposure_c \times Post_t + \alpha_t + \delta_q + \gamma_c + X_{i,c,t} + \epsilon_{i,c,t} \quad (6)$$

where $y_{i,c,t}$ is a dropout indicator for individual i in commuting zone c at time t (measured quarterly). $Exposure_c$ is an indicator equal to one if any methamphetamine laboratory was seized in commuting zone c before 2015. $Post_t$ is an indicator for periods after April 2019. The specification includes time period fixed effects (α_t), quarter-of-year fixed effects (δ_q) to account for seasonality in enrollment, and commuting zone fixed effects (γ_c). Individual-level controls $X_{i,c,t}$ include age, gender, and household education of adults. Standard errors are clustered at the commuting zone level.

The coefficient β captures the differential change in dropout rates in treated relative to control commuting zones after 2019. A negative β would indicate that exposure to synthetic drug production infrastructure led to reduced dropout rates, consistent with the theoretical prediction that skill-biased technological change in the illegal sector increases returns to education.

5.2 Tertiary Enrollment

For tertiary enrollment, the outcome is measured at the commuting zone level and includes many zeros, particularly for smaller CZs and specialized fields. I therefore estimate a Poisson pseudo-maximum likelihood (PPML) model (Silva and Tenreyro, 2006):

$$E[\tilde{Y}_{c,t}] = \exp(\beta \cdot Exposure_c \times Post_t + \gamma_c + \alpha_t) \quad (7)$$

where $\tilde{Y}_{c,t}$ is enrollment in commuting zone c in year t , allocated to origin CZs using migration shares as described in Section 4.1. PPML handles zeros naturally and provides consistent estimates under heteroskedasticity (Silva and Tenreyro, 2006). I estimate this specification separately for total enrollment, chemistry-intensive STEM+H enrollment, and non-chemistry STEM+H enrollment to test whether effects are concentrated in fields relevant to synthetic drug production.

5.3 Reaching the Final Grade of High School

The dropout specification measures contemporaneous enrollment. To measure attainment—whether students actually progress to the end of high school—I use the Formto 911 census to follow cohorts of middle-school entrants. For each commuting zone c and entry cohort τ (students entering the first grade of middle school in year τ), I compute $S_{c\tau}$, the share of the cohort that remains continuously enrolled through the final grade of high school, as the product of grade-to-grade continuation rates along the cohort’s trajectory.

Because the fentanyl ban occurred in April 2019, in the middle of the 2018–19 school year, cohorts were exposed to it at different stages of their schooling. I index each cohort by the grade $G_{c\tau}$ it was enrolled in at the ban and estimate the cohort event study

$$S_{c\tau} = \sum_{g \neq 12} \beta_g \text{Exposure}_c \cdot \mathbb{I}\{G_{c\tau} = g\} + \gamma_c + \delta_g + \varepsilon_{c\tau}, \quad (8)$$

where γ_c are commuting-zone fixed effects and δ_g are grade-at-ban fixed effects; observations are weighted by cohort size and standard errors are clustered at the commuting-zone level. The omitted category is the cohort enrolled in the final grade of high school at the ban—the last cohort to reach the terminal grade before the shock. Coefficients β_g for $g < 12$ measure the effect on cohorts that had not yet reached the final grade when the ban hit, and so completed part of their trajectory under the new regime. The coefficient of central interest is β_{11} , for the cohort enrolled in the second-to-last grade of high school (grade 11, approximately age 16) at the ban: this cohort crossed into the final grade in the first post-ban school year (2019–20), recorded in September 2019—after the April 2019 ban but before the March 2020 COVID-19 school closures—making it the one cohort whose attainment response is observed free of pandemic confounding.

I also report a pooled estimate that replaces the grade-at-ban interactions with a single $\text{Exposure}_c \times \text{Post}_\tau$ term, where Post_τ indicates cohorts that reached the final grade in or after 2019. Because every post-ban cohort except the 2019 cohort reached the terminal grade during or after the COVID-19 closures, this pooled estimate blends the disruption-free 2019 effect with COVID-affected cohorts and should be read as a conservative summary. Identification mirrors the dropout specification: absent the ban, cohort progression would have trended in parallel across successive cohorts in treated and control commuting zones, a hypothesis the pre-ban cohorts (enrolled in grades 13–15 at the ban, reaching the final grade in 2015–2017) test directly.

5.4 Identification

The key identifying assumption is that, in the absence of China’s fentanyl ban, educational outcomes would have evolved similarly in treated and control commuting zones. This assumption would be violated if pre-existing trends differed across treatment status or if other shocks coincided with the 2019 ban and differentially affected treated areas.

I assess the plausibility of parallel trends in several ways. First, I estimate event-study specifications that allow for dynamic treatment effects and visually inspect pre-trends. Second, I use the machine learning-predicted treatment measures, which address potential selection into observed seizures based on enforcement intensity.

Figure 4 presents standardized differences in baseline characteristics between treated and control commuting zones. Treated CZs differ significantly on several dimensions. In terms of geography and access, they have higher annual precipitation, greater elevation range, larger total area, and are closer to state capitals and Pacific international ports. They also show higher indicators of illegal activity, including greater cannabis and poppy suitability, more cartel presence, and larger cannabis and poppy seizures. Demographically, treated CZs are more urban (lower rural population share), have better infrastructure (higher sewerage and electricity coverage), and a larger share of working-age population. On education, treated CZs have more higher education institutions, campuses, and programs, though baseline enrollment rates and years of schooling are similar and there are no differences in the share of chemistry-intensive programs. Treated CZs also have higher labor force participation, mostly driven by more female employment, and more formal jobs. Notably, treated and control CZs do not differ in the total number of firms or the overall share of chemical industry firms. They differ only in the share of fentanyl-relevant firms—a narrower category that includes specific chemical manufacturing subsectors (such as pharmaceutical and pesticide production), medical equipment suppliers, and wholesale distributors of laboratory equipment. This pattern suggests that treated areas do not simply have larger or more industrialized economies, but rather possess the specific industrial infrastructure relevant to synthetic drug production.

These differences are consistent with drug production locating in areas with specific geographic advantages (access to ports, suitable terrain for concealment) and existing illegal economy infrastructure. While treated and control CZs differ on levels of many characteristics, the identification strategy relies on parallel trends in outcomes, not balance in levels. A natural concern is that some contemporaneous shock correlated with a single unbalanced characteristic, such as proximity to Pacific ports or cartel presence, could confound the results. The machine learning approach provides a partial response to this concern: the ML-predicted treatment is identified from the joint combination of all characteristics in Figure 4, not from any single variable. For a confounding shock to explain results under both the baseline seizure measure and the ML measure, it would need to correlate with the specific multivariate pattern that predicts synthetic drug production—a more demanding requirement than correlation with any individual characteristic.

Beyond these observable differences, several contemporaneous policy changes could confound the results if they differentially affected treated commuting zones. These include substantial minimum wage increases beginning in 2019 (with larger increases in the Northern Border Free Zone), the discontinuation of Mexico’s conditional cash transfer program Prospera, and the COVID-19 pandemic beginning in 2020. In Section 9, I present tests ruling out these alternative explanations.

6 Results

The theoretical framework in Section 3 generates a set of predictions about how skill-biased technological change in the illegal sector affects educational investment: reduced dropout and increased schooling, a reallocation of field choice toward chemistry, and a rise in the demand for chemistry skills in the legal labor market. I test each in turn.

6.1 Effects on School Dropout and Grade Progression

The model predicts that a rise in the illegal chemistry premium (π_I) raises educational investment among individuals with high propensity for illegal work. Because both sectors draw chemists from the same local pool, this bids up the legal chemistry premium (π_L) as well, so educational investment increases even among individuals with no propensity for illegal work. Both channels predict reduced dropout in treated areas. Section 7 provides complementary labor-market evidence on the rise in demand for chemistry skills that underlies these responses.

6.1.1 Event Studies and Parallel Trends

Figure 5 presents event-study estimates for dropout rates, separately for middle school age (12–14) and high school age (15–17) individuals. For both age groups, coefficients in the pre-treatment period are close to zero and statistically indistinguishable from the baseline. The mean pre-treatment estimate is -0.002 (p-value = 0.669) for middle school age and 0.005 (p-value = 0.482) for high school age, providing no evidence of differential trends prior to the announcement of the fentanyl ban.

Effects emerge after the fentanyl ban and grow over time. For middle school age, the mean post-treatment estimate is -0.012 (p-value = 0.011), indicating a 1.2 percentage point reduction in dropout rates in treated relative to control CZs. For high school age, the effect is slightly larger at -0.015 (p-value = 0.045). The gradual strengthening of effects over the post-period—with estimates becoming more negative by year 5—aligns with the theoretical prediction that educational responses take time to materialize as younger cohorts adjust their investments.

However, the timing of effects must be interpreted with caution due to the COVID-19 pandemic. Beginning in March 2020—just one year after the fentanyl ban—Mexico closed schools and transitioned to distance learning. Schools remained closed until the beginning of the 2021–2022 academic year (Q4 2021, 1.5 years after the fentanyl ban). During the closure period, students were automatically re-enrolled and promotion policies were relaxed, making it difficult to observe dropout in the administrative sense. This implies that even if the fentanyl ban had immediate effects on educational decisions, they would be masked in the data during 2020–2021. The effects I observe emerging in years 2–5 post-shock may therefore reflect both the true gradual adjustment predicted by the model and the mechanical suppression of dropout during school closures.

As a placebo test, I also examine effects on primary school age children (6–11). Figure 6 shows no evidence of treatment effects for this age group. This null result is consistent with the human capital investment interpretation: children of primary school age are not at the margin of dropping out to enter the labor market, so changes in skill returns should not affect their enrollment. The concentration of effects among middle and high school age youth—who face meaningful tradeoffs between continued schooling and employment—supports the view that reduced dropout reflects responses to changing labor market returns rather than other shocks that would affect all age groups equally, such as changes in school quality or household income.

6.1.2 Main Estimates

Table 1 presents difference-in-differences estimates of the effect of exposure to synthetic drug production infrastructure on dropout rates. Columns (1)–(4) report results for middle school age and columns (5)–(8) for high school age. All specifications include commuting zone, year-quarter, and quarter fixed effects, with standard errors clustered at the commuting zone level.

For middle school age, the baseline estimate in column (1) indicates that exposure to synthetic drug production reduces dropout by 1.7 percentage points ($p < 0.01$), a 29% reduction relative to the mean dropout rate of 5.9%. This estimate is stable across specifications that progressively add individual controls for gender and age (column 2), household average education of adults (column 3), and an interaction with gender (column 4). The coefficient on female is negative and significant, indicating that girls have lower dropout rates, while older children and those in households with less educated adults are more likely to drop out. The interaction between treatment and female in column (4) is close to zero and insignificant, suggesting similar effects for boys and girls.

For high school age, effects are larger in magnitude. The baseline estimate in column (5) shows a 2.1 percentage point reduction in dropout ($p < 0.05$), representing an 8.6% reduction relative to the higher baseline dropout rate of 24.5%. Effects remain robust, ranging from 2.2 to 2.7 percentage points across specifications (columns 6–8). As with middle school, female students have substantially lower dropout rates, older students are more likely to drop out, and household education is protective. The interaction with gender in column (8) is again small and insignificant. The larger absolute effects for high school age are consistent with this group being closer to the margin of labor market entry, where changes in skill returns are most salient.

6.1.3 Effects by Quarter

To further understand when educational decisions are made, I estimate the main specification separately for each quarter. Table 2 presents these results. In Mexico, the school year runs from August to July, with Q4 (October–December) corresponding to the first quarter of the academic year. If effects were driven solely by enrollment decisions at the start of the school year, we would expect effects to be concentrated in Q4.

Overall, the difference-in-difference estimates by quarter in Table 2 are consistent with the pooled estimates. For middle school age, the estimates are similar across quarters, ranging from -0.012 to -0.018 , suggesting that exposure to synthetic drug production reduces dropout relatively uniformly throughout the school year for this age group. As before, the event study evidence in Figure 7 shows effects that strengthen over the post-period. Evidence for Q1 is nonetheless particularly informative for timing. Since the last pre-ban Q1 (January–March 2019) serves as the reference period, the first post-treatment Q1 observation corresponds to January–March 2020—exactly one year of exposure to the fentanyl ban, and the maximum duration that can be observed before COVID-19 school closures disrupted enrollment measurement starting mid-March 2020. This estimate is already negative and statistically significant, demonstrating that short-term effects could have emerged rapidly following the policy shock.

For high school age, the pattern is more striking. The magnitude increases throughout the school year: from -0.014 in Q4 (start of year, not statistically significant) to -0.033 in Q2 (end of year). The Q2 effect represents a 12.5% reduction relative to the quarter-specific mean dropout rate of 26.3%. This concentration of effects at the end of the academic year suggests that high school students who would have otherwise dropped out before completing the year are instead persisting through graduation, consistent with revised expectations about returns to completing their education.

The event study evidence by quarter reinforces these findings. The end-of-year quarter (Q2) shows the clearest pattern (Figure 7): flat pre-trends followed by effects that grow steadily larger over time, with the most pronounced reductions in dropout occurring in the later post-treatment years. Effects in the start-of-year quarter (Q4) are negative but smaller and less precisely estimated.

6.1.4 Time Use Evidence

A potential concern with the dropout results is that COVID-19 school closures beginning in March 2020 mechanically froze enrollment patterns, making it difficult to observe true educational investment during 2020–2021. To address this, I examine time use data from the same survey (ENOE), which captures how school-age individuals allocate their time across activities regardless of formal enrollment status.

Figure 8 presents event study estimates for weekly time spent studying among middle and high school age youth, estimated by Poisson regression so that coefficients are proportional effects (log points). The pre-treatment coefficients show no differential trend (mean estimate 0.016, $p = 0.22$). Following the fentanyl ban, time spent studying increases sharply in treated areas (mean post-treatment estimate 0.053, roughly a 5% increase, $p = 0.04$), with effects that persist throughout the post-period—including during the COVID-19 school closures when formal enrollment was administratively frozen.

Table 3 presents difference-in-differences estimates for various time-use categories; the indicator for

any studying is estimated by a linear probability model and the remaining (minutes) outcomes by Poisson, so their coefficients are proportional effects. Exposure to synthetic drug production raises the probability of any studying by 2.1 percentage points for middle school age and 3.3 percentage points for high school age (both $p < 0.001$), and increases time spent studying by roughly 2% ($p < 0.01$) and 5% ($p < 0.001$), respectively. Time spent on household chores falls by about 5% for middle school age and 6% for high school age (both $p < 0.001$), while time spent on care work and shopping shows no significant change. This pattern suggests that increased educational investment comes partly at the expense of household production, consistent with a reallocation of time toward human capital accumulation rather than a general increase in productive activities.

6.1.5 Reaching the Final Grade of High School

The dropout and time-use results show that exposure raises contemporaneous enrollment and study effort. A distinct question is whether this translates into greater educational *attainment*: the ENOE records whether a student is attending school at the survey date, not whether the student ultimately advances through the system. I therefore turn to the Formato 911 census, which tracks grade-to-grade progression for the universe of schools and allows me to measure whether cohorts of students reach the final grade of high school.

Figure 9 presents the cohort event study. Cohorts that reached the final grade of high school before the ban (in green) lie on a flat, statistically insignificant path, indicating no differential pre-trend in cohort progression between treated and control commuting zones. The effect appears sharply for the first exposed cohort. Students enrolled in the second-to-last grade of high school (grade 11, approximately age 16) when the ban was announced were 6.2 percentage points more likely to reach the final grade in treated relative to control zones ($p < 0.001$)—an increase of about 11% relative to the control-group rate of 54.9%. This cohort crossed into the final grade in September 2019, after the ban but before the March 2020 school closures, so its response is measured free of pandemic disruption.

Subsequent cohorts, which reached the final grade during or after the COVID-19 closures, show positive but smaller and noisier effects (Appendix Table A.1); the pooled post-ban estimate is 3.1 percentage points ($p < 0.01$), attenuated by the inclusion of these disrupted cohorts. That the pandemic *dampens* rather than amplifies the estimate—the cleanest, pre-COVID cohort exhibits the largest effect—indicates that the school-closure confound works against, not in favor of, the finding.

This result disciplines the interpretation of the dropout estimates in two ways. First, it is an *attainment* effect: it measures whether cohorts advance to the terminal grade of high school, not merely whether they are enrolled at a point in time. Second, because it is identified from a single pre-COVID cohort, it addresses the concern that the ENOE dropout effects are inflated by pandemic-era enrollment freezes. Taken together, the survey and administrative evidence indicate

that the contemporaneous rise in enrollment translates into a genuine increase in the share of students who reach the final grade of high school.

6.2 Effects on Field of Study

The field-choice model delivers a sharp *compositional* prediction: because the shock raises the chemistry premium, students reallocate toward chemistry-intensive programs, so chemistry enrollment rises while non-chemistry and total enrollment change only through the (small) extensive margin. The tertiary-enrollment results test this prediction directly.

6.2.1 Event Studies and Parallel Trends

Figure 10 presents event study estimates for tertiary enrollment by field of study. For overall enrollment, there is some evidence of a differential pre-trend: the coefficients are positive but decline toward zero in the years just before the fentanyl ban (mean pre-treatment estimate 0.028, $p = 0.02$). This suggests that treated commuting zones were on a relatively declining tertiary-enrollment path prior to 2019, potentially reflecting broader demographic or economic trends in areas with drug production infrastructure. The post-treatment coefficients are close to zero early on but increase steadily over time.

For chemistry-intensive STEM+H programs—including chemical engineering, pharmaceutical sciences, and biotechnology—the picture is strikingly different. The pre-treatment coefficients show no evidence of differential trends, with estimates clustered around zero throughout the pre-period. Following the 2019 fentanyl ban, effects emerge clearly and grow steadily over time. By year 4 after the shock, the coefficient reaches approximately 0.09, corresponding to an $e^{0.09} - 1 \approx 9\%$ increase in chemistry-related enrollment in exposed commuting zones relative to control areas. This strengthening pattern is consistent with cohorts who were exposed to the new labor market conditions during secondary school progressively reaching university age and selecting into chemistry-related fields.

For non-chemistry STEM+H programs, the event study reveals no discernible pattern—neither pre-trends nor post-treatment effects.

6.2.2 Main Estimates

Table 4 presents the corresponding difference-in-differences estimates using Poisson pseudo-maximum likelihood, which naturally handles the many zeros in field-specific enrollment. Panel A shows results for all students. The estimate for chemistry-intensive STEM+H programs is 0.057 ($p < 0.05$), corresponding to a $(e^{0.057} - 1) \times 100 \approx 5.9\%$ increase in enrollment in exposed areas after the fentanyl ban. The estimates for non-chemistry STEM+H and total enrollment are small and statistically insignificant. This is precisely the compositional pattern predicted by the field-choice model: exposure raises chemistry-intensive enrollment while leaving non-chemistry and total en-

rollment flat, consistent with students reallocating *across* fields rather than a uniform increase in tertiary attainment.

Panels B and C decompose these effects by gender, revealing different margins of adjustment for women and men. Both groups show increased enrollment in chemistry-intensive programs: women by 0.061 log points ($p < 0.01$), or approximately 6.3%, and men by 0.048 log points ($p < 0.10$), or approximately 4.9%. However, the source of this reallocation appears to differ by gender.

For women, the increase in chemistry-intensive enrollment is accompanied by a marginally significant decrease in non-chemistry STEM+H enrollment (-0.049 , $p < 0.10$), with no change in overall enrollment. This pattern suggests reallocation within STEM+H fields: women who would have pursued other technical degrees—such as non-chemical engineering or health sciences—are instead selecting into chemistry-related programs.

For men, the picture is different. The increase in chemistry-intensive enrollment is not offset by declines in non-chemistry STEM+H, which remains unchanged. This suggests that men entering chemistry-related fields are drawn either from non-technical fields or from the margin of not attending university at all, rather than substituting away from other STEM+H programs.

These gendered patterns are consistent with different baseline constraints. Women in STEM+H may face a relatively elastic choice among technical fields, responding to shifts in skill-specific returns by reallocating within a set of programs they were already considering. Men’s response, by contrast, may operate on a more extensive margin—drawing individuals into chemistry-intensive tertiary education who would not otherwise have pursued technical training.

This pattern of results provides strong support for the field-choice mechanism: the technological shift toward synthetic drug production induced students to reallocate toward fields providing chemistry-specific skills. The contrast between the null effects on overall and non-chemistry enrollment and the positive, growing effects on chemistry-intensive programs rules out alternative explanations based on general improvements in educational infrastructure, broad economic development in treated areas, or increases in STEM interest unrelated to the specific skill demands of synthetic drug production. The gender differences further suggest that women may be more responsive to shifts in skill-specific returns, potentially because they face different baseline constraints or opportunities in the labor market.

7 Mechanisms

The educational and field-of-study responses documented in Section 6 rest on a premise: that exposure to synthetic-drug production raised the demand for chemistry skills in a way individuals could act on. Two features of the setting must be established before that demand can be read off wages and employment. Section 7.1 shows that exposure did not bring a violence shock, so the risk of entering the illegal sector did not rise and a change in violence is not itself the driver of the

schooling response. Section 7.2 establishes what the ENOE observes—legal employment, formal and informal—and shows that the workers drawn into concealed production drop out of the sample. With these premises in place, the price, quantity, and composition of chemistry employment in the sections that follow can be read off the legal sector.

7.1 Low-enforcement and unchanged violence environment

Exposure to synthetic-drug production did not raise local violence, a fact that matters for two reasons unrelated to any labor-market return: whether a change in violence is itself behind the schooling response, and whether entering illegal production was dangerous enough to make the response implausible.

A decline in violence is one candidate explanation for the schooling results, and the data rule it out. Violence is a well-documented suppressant of human-capital accumulation in Mexico (Jarillo et al., 2016; Brown and Velásquez, 2017), so a fall in local violence could by itself raise educational investment—though it could not explain why that investment flows specifically into chemistry. The homicide evidence below shows no such decline, which closes this channel.

Entering illegal production in this setting did not carry the near-certain detention or death that the term “drug production” evokes, and this is central to the plausibility of the response. A reader might reasonably doubt that people would move toward an activity that seemed to promise prison or a violent end; in Mexico during the outcome window, it did not. Federal enforcement of fentanyl production was minimal for most of the period: clandestine laboratories were seized, but the people operating them were very rarely detained, and there were no major fentanyl-related arrests until late in the sample. High-value captures did not fragment the organization—Ovidio Guzmán was briefly held in October 2019 and again in January 2023, yet production continued in both cases (Dell, 2015)—and the United States designated the Sinaloa Cartel under counter-narcotics sanctions only in December 2021 (U.S. Department of the Treasury, 2025). The first structural disruption, the 2024 split of the cartel, post-dates the outcome window (Sobrino, 2019). Entry into production was thus a bet on earnings, not a gamble with liberty or life, which is what makes a return-driven response credible.

The homicide data confirm that violence did not rise. Figure 13 presents the event study for the intentional-homicide rate, estimated by Poisson regression with a log-population offset—using annual commuting-zone population from CONAPO’s municipal projections—so that coefficients are proportional changes in the per-capita rate. Treated and control zones share a flat pre-trend (mean pre-treatment estimate -0.04 , $p = 0.66$), and the rate does not change significantly after the ban (post-treatment mean 0.10 , $p = 0.24$; the pooled estimate corresponds to a 14% change, $p = 0.18$). If anything the estimates are slightly positive—the opposite of a decline—and among the commuting zones already most violent before the ban the increase is larger and marginally significant (+21%, $p = 0.07$), still far below the scale of a genuine violence shock. The result holds

using population fixed to the five-year census anchors and using the unweighted homicide count (Appendix Table G.1 and Appendix Figure G.1).

Whatever change in homicides does occur is too gradual and too late to drive the results. The post-treatment path starts near zero at the ban (+0.03 at event time zero) and drifts up only over subsequent years, whereas the educational responses of Section 6 appear immediately. A slow-moving, marginal change in violence cannot account for effects that turn on at impact.

7.2 Shifting employment towards illegal activities observed by increased concealment in the labor-market survey

The labor force survey, ENOE, observes what is most likely legal employment—formal and informal alike—but not concealed illegal production, and this distinction governs how the labor-market results are read. Its coverage of informal work is exactly what suits it to the question: because more than half of Mexican employment is informal, a survey confined to registered jobs would miss most of the labor market, and because the ENOE is a household survey it records workers whether or not their employer is registered. What it cannot record is clandestine drug production. No prior work establishes how the survey behaves in areas of illegal production, so I build the case here from its structure, the pattern of non-response, and what can be ruled out.

Because illegal production cannot be reported as employment, the workers engaged in it either stay in the sample—appearing as legally employed or out of the labor force—or drop out of it, and the non-response evidence points to the latter. That the ENOE already counts informal work sharpens the inference: a worker who disappears from it is not merely informal, since informal jobs are recorded, but is concealing activity the survey would otherwise capture.

Survey refusals rise sharply with exposure, turning on precisely at the ban. Figure 14 presents the event study for the household refusal rate—the share of inhabited dwellings whose occupants decline to be interviewed. Treated and control commuting zones are on a common pre-trend (mean pre-treatment estimate -0.003 , $p = 0.31$), and refusals then rise over event time to a post-treatment mean of 1.2 percentage points ($p = 0.009$), a substantial movement against a control-group refusal rate of 2.7%; the pooled difference-in-differences estimate is 0.7 percentage points ($p < 0.05$, Appendix Table G.2).

The rise is not out-migration. If residents had left, the dwelling would be recorded as uninhabited or vacated—a non-interview reason the ENOE codes separately from refusal by present occupants—so departures do not enter this margin. The increase is in refusals by households that are there, not in empty dwellings, and migration itself is examined directly in Section 7.4.

The rise is not fear of victimization. Households might decline a stranger at the door if they felt less safe, but the incentives of the shock run the other way: higher fentanyl rents raise the opportunity cost of preying on the local population, so organizations profiting from production have less reason

to extort or victimize residents. No municipal extortion series lets me test this directly, but it is consistent with the homicide evidence of Section 7.1—there is no violence shock for households to be reacting to.

The rise is not evasiveness within completed interviews. If exposed households were merely more guarded with the interviewer, it would surface as item non-response among those who answer—yet it does not (Appendix Table G.2, Panel B). Missing labor income does not rise; occupation is reported as completely as before; and income *heaping*—the tendency to round reported earnings to multiples of 1,000 or 5,000, a standard proxy for imprecise reporting—if anything falls, so incomes are reported slightly more precisely, not less. That respondents remain forthcoming even on earnings, the most sensitive question, is a further strike against the victimization reading: households that had grown wary of outsiders would answer less fully, not more. The margin that moves is entirely at the unit level—whole-household refusal—not selective evasion inside interviews that are completed.

The additional refusals therefore come, most plausibly, from households connected to concealed production, who decline the survey rather than misreport within it. This reading is consistent with accounts from ENOE field staff at INEGI, who describe recurring difficulty completing interviews in areas of known illicit production.

Two implications carry forward. First, the observed sample is disproportionately the legal labor market: the workers most likely to be in concealed production are precisely those now missing from it, so the chemistry wage premium, the employment expansion, and the compositional shift documented below are all measured on legal, observed employment, with concealed workers selectively absent—and where their absence would move an estimate in a known direction, I flag it. Second, the exit is selective, not wholesale: a 1.2-percentage-point rise off a low base is a shift at the margin, not a survey gone blind to the treated economy, so it is enough to sign the interpretations that follow without treating the ENOE as though it observed only legal activity.

7.3 Demand for chemistry skills rising in the legal labor market

The shock raised both the price and the quantity of chemistry skill in the legal labor market—the joint movement that identifies a demand shift. Table 6 reports it: Panel A the return to a chemistry-relevant degree (the price), Panel B employment in chemistry-relevant occupations (the quantity), each estimated with the paper’s difference-in-differences design on prime-age (25–54) workers.

The return to a chemistry degree rose, while the return to schooling in general did not. The wage premium on a chemistry-course field increased by 0.059 log points in exposed commuting zones after the shock (Panel A, column (6), $p < 0.05$), against a baseline premium of 0.52; the return to STEM as a whole rose by a similar 0.050 ($p < 0.05$), but this is concentrated in its chemistry component—the return to STEM fields without chemistry content is half as large and insignificant. The overall wage effect and the return to non-STEM and non-chemistry degrees are all flat. Because

Section 7.2 established that the ENOE measures the legal market, this is the legal price of chemistry skill, measured cleanly rather than blended with illicit pay. It is also a lower bound: the highest-return chemists—those drawn into concealed production—are the ones selectively missing from the sample, so their departure depresses the measured premium. And it cannot come from front hiring, which pays ordinary wages for ordinary work; a *rising* price of chemistry skill is evidence of genuine demand for that skill.

Employment in chemistry occupations rose at the same time. The share of employment in chemistry occupations increased by 0.176 log points (Panel B, column (6), $p < 0.01$), roughly a 19% gain, and employment in STEM occupations more broadly by 0.073 ($p < 0.01$); the overall employment rate and non-STEM professional employment did not move. As with degrees, the expansion is specific to chemistry-intensive technical work, not a general rise in employment. Figure 15 traces the dynamics: the STEM-degree wage premium and chemistry-occupation employment are flat before the ban and rise after it, with no differential pre-trend.

Price and quantity rising together is the signature of an outward demand shift, and it is a signature the alternatives cannot reproduce. A pure inflow of chemistry-skilled workers—labor supply expanding down a fixed demand curve—would raise employment but push the wage *down*; poaching of scarce chemists would bid the wage up while employment fell. Only an outward shift in the demand for chemistry skill raises the price and the quantity at once. This is the discriminating test that a violence- or displacement-based account fails: those stories move quantities without raising the price of the specific skill. What the legal labor market reveals, then, is exactly the change in chemistry-skill demand that the educational responses of Section 6 were responding to.

7.4 Local, less-educated workers are filling the chemistry jobs in the legal sector

The expanding chemistry workforce is filled by drawing in less-credentialed local workers, not by credentialed chemists—a reading that the education and the geographic origin of that workforce both support. Table 7 documents both, using the same commuting-zone $\times$ year design and occupation columns as the employment table.

The workforce of chemistry and STEM occupations becomes measurably less educated. Panel A reports the average years of schooling of each occupation group’s workers: it falls by 0.15 years in STEM occupations ($p < 0.01$) and by 0.28 years in chemistry occupations ($p < 0.05$), while schooling in professional and non-STEM work does not move. This is the expected companion to the employment expansion of Section 7.3, not an anomaly. Expanding an occupation’s employment moves the market out along an upward-sloping labor-supply curve and draws in the marginal worker, who by construction sits below the incumbent average, so the group’s mean schooling falls mechanically. A rising price, a rising quantity, and a marginal entrant below the incumbent average is the textbook signature of a demand shift, and it needs no new assumption. It also leaves the premium of Section 7.3 conservative: because the entrants are less credentialed and therefore

lower-paid, their arrival pulls the measured premium down rather than up.

The drawn-in workers plausibly carry chemistry competencies acquired below the degree level, and some of the measured growth may fall in ancillary support roles; I offer this as interpretation, not as something the data test.

The new employment is locally sourced. Panel B reports the out-of-state (migrant) share of each group’s workforce: it falls for the employed as a whole (-0.9 percentage points, $p < 0.1$) and is flat and insignificant within STEM and chemistry occupations specifically. No inflow of chemists arrives to staff the new jobs; the additional employment is drawn from the local population—which is where the schooling response of Section 6 supplies it. The same fact points forward: a *falling* out-of-state share is the opposite of what skilled out-migration would leave behind, since departing credentialed locals would tilt the remaining workforce *more* toward out-of-state origin, not less. Exposed areas are retaining their local chemistry workers, not losing skilled people. This origin mix concerns employed chemistry workers, a different population from the degree-holders examined next, so it corroborates rather than settles the retention reading of Section 7.5.

7.5 Individuals with chemistry credentials are being poached into illegal production

Among people whose schooling was already complete before the shock, the local share holding a chemistry or STEM degree falls in exposed commuting zones—and the fall is best read as credentialed workers leaving the observed economy, not the region. This is the result that ties the mechanism together.

The decline is field-specific and predetermined. Table 8 tracks a fixed cohort—those aged 30–54 in 2019, followed across survey years by birth cohort—so their attainment cannot respond to the shock. In exposed areas the cohort’s STEM-degree share falls by 3.7% ($p < 0.1$) and its chemistry-degree share by 4.6% ($p < 0.05$), while the placebos, any tertiary degree and non-STEM field degrees, stay flat. Figure 16 shows the STEM-degree decline has no pre-trend and emerges at the ban. Because attainment is fixed, this cannot be a schooling response; it is mobility or observability. And because the placebos do not move, it is not generic educated out-migration or survey non-response, which would thin every field alike—it is specific to the fields the shock rewards.

That the credentialed *stay* in the region follows from the paper’s own logic and needs no new data. The return to chemistry skill rose in the treated commuting zone itself—in the legal market (Section 7.3) and, by the model, by more still in illicit production (Section 3). A chemist chasing that return has no reason to leave: the opportunity is where she already is. Return-chasing predicts retention, and its highest-return form—capturing the illicit rent—means entering concealed production *inside* the commuting zone, that is, exiting the observed frame while staying put. The very logic that drives the paper, that chemists follow the chemistry return, therefore predicts stay-and-go-dark, not emigrate.

The two emigration channels are closed by results already in hand. The push channel, flight from violence, is ruled out because homicides do not rise with exposure (Section 7.1): there is no violence shock to flee. The pull channel, leaving for a better return elsewhere, is ruled out because the return rose *here*; the pull is inward, not outward, with nothing to chase in another region.

The decisive argument is internal consistency. A physical exodus of the skilled is very hard to square with chemistry employment *rising* (Section 7.3). Skilled workers exiting into concealment within the commuting zone, their observed jobs backfilled by less-credentialed locals (Section 7.4), is the only reading under which all four results cohere at once: the price of chemistry skill rises, observed chemistry employment rises, the observed workforce becomes less credentialed, and the observed credentialed stock shrinks. Emigration breaks that coherence—it cannot raise local chemistry employment while draining the area of credentialed chemists—whereas concealment preserves it. This is the argument that does the work; the migration evidence of Section 7.4, where the out-of-state share falls rather than rises, confirms it rather than carrying it.

7.6 Summary

Taken together, the labor-market evidence describes a single process. Exposure to synthetic-drug production raised the legal price of chemistry skill and expanded chemistry employment at the same time (Section 7.3)—the joint movement that identifies an outward shift in the demand for chemistry skill rather than poaching, an elastic supply inflow, or violence-driven displacement. That expansion drew in less-credentialed workers (Section 7.4), sourced from the local population rather than an influx of chemists (Section 7.4), while the observed stock of credentialed chemists in exposed areas fell—leaving the observed economy, not the region (Section 7.5). The local schooling response of Section 6 is where the new supply originates, and the low-enforcement, low-violence setting (Section 7.1) is what lets the response take the form of human-capital investment rather than flight. One identification note runs through all of it: the difference-in-differences design establishes that these outcomes moved together in exposed commuting zones; that the missing chemists are concealed *within* the region rather than gone from it is an interpretation—supported by the field-relative placebos, the retained-locals migration pattern, and the absence of a violence shock—not something the research design settles on its own.

8 Size of the illicit sector

The difference-in-differences design delivers two effects that a labor market participant could in principle observe: a rise in the legal-market chemistry premium and an increase in chemistry-related schooling. It does not deliver their illegal-sector counterparts—by how much the *illegal* chemistry premium rose, or how large the demand shock that drove the episode was. This section uses the model of Section 3 to recover those two hidden objects from the observed legal premium. Because both sectors hire chemists from a single local pool, the market-clearing conditions can be

read in reverse: a legal premium of a given size can only have arisen from an illegal premium and a demand shift of particular magnitudes, and inverting the conditions recovers them.

The inversion turns on two quantities that the reduced-form estimates leave open: the *reallocation elasticity* κ_L , which governs how freely chemists move into legal chemistry work—from illegal work or from other legal occupations—and the *relative sector size* r , the illegal chemist pool as a fraction of the legal one. I therefore report the recovered magnitudes as functions of these two parameters (Figure 17) rather than as point estimates, and confront them with independent evidence on the illegal sector’s size and growth. The observed responses fit the skill-demand mechanism comfortably: rationalizing the estimated legal premium requires the illegal chemistry premium to have risen by about a quarter and illegal demand for chemistry skill to have grown several-fold to roughly thirtyfold—magnitudes that a small illegal sector facing a large proportional shock produces with room to spare, and that are, if anything, modest beside the independent growth in clandestine-laboratory and fentanyl seizures. Throughout, the exercise is a plausibility and consistency check on the mechanism, not a source of identification; identification rests on the difference-in-differences design.

Inputs. The difference-in-differences estimates give $\widehat{d \ln \pi_L} = 0.0589$ log points for the legal chemistry premium; reductions in middle- and high-school dropout of 1.7 and 2.1 percentage points, on control means of 5.9 and 24.5 percent; and a PPML coefficient of 0.057 on enrollment in chemistry-related tertiary degrees—a 5.9 percent increase—with no effect on overall tertiary enrollment, confirming the compositional content of Prediction 3.

The inversion. The shock enters only the illegal market, so the legal premium can rise only through cross-sector poaching: a higher illegal premium pulls chemists across the sorting margin, legal supply contracts, and π_L rises to re-clear the legal market. Unwinding the observed $\widehat{d \ln \pi_L}$ therefore proceeds in two steps, each using one clearing condition: the legal-market condition recovers the implied *illegal* premium $d \ln \pi_I$, and the illegal-market condition then recovers the implied demand shift $d \ln A$.¹⁶ The economics is transparent. For π_L to rise by 5.9 log points, the legal market must give up roughly $(\sigma_L + \alpha_L) \times 5.9 \approx 6.5$ percent of its chemists—the movement along legal demand plus the offset from its own weak supply response. That absolute outflow, absorbed by a sector of relative size r , is a proportional expansion of $6.5/100r$ log points of the illegal workforce, which illegal demand growth must finance.

Two features of the inversion organize what follows and shape the figure. The implied illegal premium depends on κ_L but is exactly invariant to r : it is fixed entirely by how hard the legal market made it to extract the observed flow of chemists. The implied demand shift depends on

¹⁶From the legal-market condition, using the identity $-\varepsilon_{LI} = \kappa_L \lambda$, the pass-through ψ gives $d \ln \pi_I = \widehat{d \ln \pi_L} / \psi = [(\sigma_L + \alpha_L + \kappa_L) / (\kappa_L \lambda)] \widehat{d \ln \pi_L}$, where $\lambda \equiv (1 - p(\bar{\theta}_C)) \pi_I / \pi_L = 1 - R(\bar{\theta}_C) / \pi_L \in (0, 1]$ (Case A). From the illegal-market condition, $d \ln A$ follows exactly; because $\kappa_I = \kappa_L \lambda / r$ is large when the illegal sector is small, it is well approximated by $d \ln A \approx [(\sigma_L + \alpha_L) / r] \widehat{d \ln \pi_L}$ plus order-one terms.

r but only negligibly on κ_L : the flow the legal market surrenders is pinned by the legal-market response, and r only rescales it into a proportional illegal-sector expansion. Each recovered object thus maps to a single unknown, and Figure 17 plots the two relationships in turn—Panel (a) the illegal premium against κ_L , Panel (b) the demand shift against r .

Parameters. The parameters fall into three grades of credibility. *Disciplined by the paper’s own estimates:* the creation elasticity α_L —the response of the *stock* of legal chemists to the premium over the study window—follows from the PPML estimate. The enrollment *flow* responds with elasticity $0.057/0.0589 \approx 0.97$, but with a working life of roughly 35 cohorts and only two to five treated cohorts entering the workforce by the end of the window, the stock elasticity is $\alpha_L \approx 0.97 \times (2-5)/35 \in [0.06, 0.14]$; I use 0.10. This short-run inelasticity, at the heart of the amplification, is implied by the data rather than assumed. *Anchored by external evidence:* the legal demand elasticity σ_L is set to 1, the upper end of estimated own-wage labor demand elasticities (which cluster between 0.3 and 1); because the implied magnitudes increase in σ_L , this is conservative, and the figure traces $\sigma_L \in [0.75, 1.5]$. The illegal-side elasticities ($\sigma_I + \alpha_I = 0.85$) are quantitatively irrelevant next to $\kappa_I \in [8, 60]$. *Without direct benchmarks:* the remaining two parameters have no clean empirical counterpart, so rather than fix a point value for either I let Figure 17 trace the implied magnitudes across their full plausible ranges. The reallocation elasticity κ_L —the percent of legal chemist supply gained per percent rise in π_L as chemists move into legal chemistry work, whether drawn back from illegal work or in from other legal occupations—has no direct estimate in the literature, and Panel (a) spans an order of magnitude, from sluggish (0.2) to fluid (1.5) reallocation.¹⁷ The relative sector size r —clandestine synthesis personnel with chemistry skills against legal chemistry-related employment—is bounded below by measurement (discussed under external benchmarks), and Panel (b) spans 1 to 5 percent. Finally, $\lambda = 1 - R(\bar{\theta}_C)/\pi_L \in (0, 1]$; I use 0.8, and since it enters only through $\kappa_L \lambda$, any misallocation is absorbed by the κ_L axis of Panel (a).

Implied magnitudes. Panel (a) of Figure 17 shows that rationalizing the observed 5.9 log-point legal premium requires the illegal chemistry premium to have risen between 14 and 61 percent, with +27% at the central reallocation elasticity. None of these values is implausible: qualitative accounts describe extraordinary pay for cartel-recruited synthesis chemists, so the mechanism does not require premiums outside the range of reported behavior even under sluggish reallocation. And precisely because the whole range is admissible, the premium estimate alone cannot pin down the degree of labor-market integration; what disciplines κ_L is the *quantity* evidence—the growth of chemistry-related employment in the legal economy of exposed areas (Section 7), which shows the reallocation margin is active and argues against the sluggish end of the panel, and with it against the largest implied illegal premiums. Panel (b) converts the question “is the illicit sector big enough?”

¹⁷Reallocation cannot quickly create genuinely new synthesis-capable chemists, so the short-run inelasticity applies to the *skilled* pool that commands the premium. Less-credentialed workers drawn into ancillary chemistry-related roles raise chemistry-occupation *employment*—the de-skilling of Section 7—without relieving the scarcity of genuine synthesis skill.

into a measurement question. Given the legal-market elasticities, the legal economy surrendered roughly 6.5 percent of its chemists; how large a proportional demand boom that implies depends only on the base it landed on. If clandestine production employed on the order of 2 percent as many chemistry-skilled workers as the legal sector, illegal demand for chemistry skill grew about thirtyfold over 2019–2024; at 5 percent, several-fold. A *small* sector experiencing a *large* proportional shock, facing skill supply that the paper’s own PPML estimate shows to be inelastic, can therefore move the premium and hence education outcomes.

External benchmarks. The mechanism does not require the illegal sector to be large, but it does require it to be *non-negligible*: were clandestine chemistry to employ far below one percent as many workers as the legal sector, reproducing the observed legal premium would demand a shock too large to credit (Figure 17, Panel b). The relevant question is therefore not whether the illegal sector is small—it plainly is—but whether it is nonetheless large enough to move the market. Three independent pieces of evidence indicate that it is, placing r at roughly one to five percent. First, the *level* of chemistry employment. In the exposed commuting zones the average zone employs about 675 workers in chemistry-related occupations, some 92,000 in all.¹⁸ For r to reach one to five percent, cartel chemistry need employ only about 7 to 34 chemists per zone—some 900 to 4,600 in all—a floor that hundreds of clandestine laboratories, each staffed by a few skilled cooks, clear comfortably: there are easily enough illegal chemists to make up the required share. Second, the *scale* of criminal employment. Prieto-Curiel et al. (2023) estimate that Mexican cartels employed 160,000 to 185,000 people in 2022, about 0.3 percent of the national workforce and a floor for the exposed zones, where production concentrates; against a criminal workforce this large, a chemist ratio of a few percent is unremarkable. Third, the *growth* of production. Over the study window—matching the difference-in-differences design, comparing the post-ban to the pre-ban years and dropping the 2019 transition year—detected clandestine-laboratory seizures grew about 4.4-fold and fentanyl seizures about 8.7-fold, and both are lower bounds, since weak enforcement (Section 7.1) leaves much production undetected. Read through the inversion, a demand shock of this size reproduces the observed premium at r in the low single digits (Figure 17, Panel b), consistent with the level evidence. The three readings coincide: whether measured from the size of the legal chemistry pool, the scale of criminal employment, or the growth of production, r lands in the low single digits—small in absolute terms, but an order of magnitude too large for the “too small to matter” objection to hold.

Internal consistency. A final check compares the education and premium responses directly. The ratio of the schooling response to the premium response identifies the density of marginal students scaled by the premium: the estimates imply elasticities of school completion with respect

¹⁸ Author’s calculations from the Mexican labor force survey (ENOE), 2015–2018, exposed commuting zones, survey-weighted; chemistry-related occupations comprise laboratory, chemical, pharmaceutical, and agro-industrial occupations and chemistry teaching. The estimate excludes Mexico City, whose commuting zone is a large outlier (about 90,000 such workers); including it raises the average to about 1,325 per zone.

to the chemistry premium of 0.47 (high school) and 0.31 (middle school), and an elasticity of chemistry-degree enrollment of $0.057/0.0589 \approx 0.97$. All three are positive, ordered sensibly (largest at the margin closest to the field choice, smallest where baseline completion is nearly universal), and within the range of schooling responses to returns documented in the literature.¹⁹

Caveats. Three caveats discipline the exercise. First, it is a plausibility and consistency device: the calculations show the observed premium and education responses are *consistent with* the skill-demand mechanism under defensible elasticities; they do not identify it, and no causal claim rests on them. Identification comes from the difference-in-differences design of Section 6. Second, the check has content because it could have failed: had the education responses been an order of magnitude larger relative to the premium, the implied density of marginal students would have been indefensible, and had the premium required an implied illegal premium or demand shift outside any plausible range, the mechanism would have been strained; the estimates instead sit comfortably in the admissible region. Third, the calculations condition on Case A and an active sorting margin, which the labor-market evidence supports but does not prove.

9 Robustness

9.1 Robustness to Alternative Treatment Definitions

A key concern with the baseline treatment measure—any methamphetamine laboratory seized before 2015—is that seizures reflect both underlying production capacity and enforcement effort. Areas with intensive enforcement may have more seizures despite similar production levels, while areas with weak enforcement may have undetected production. To address this, I use machine learning methods to predict latent production capacity based on geographic, economic, and demographic characteristics, independent of enforcement intensity.

I train Random Forest, XGBoost, and ensemble models to predict the probability of synthetic drug production using features measured in 2015 or earlier, including distance to U.S. border and Pacific ports, proximity to universities with chemistry programs, presence of chemical firms, cartel activity, and geographic characteristics favorable to concealment. Crucially, the models exclude enforcement-related variables. I implement spatial cross-validation, partitioning Mexico into five geographic folds and training on four while predicting on the held-out fold, which forces the model to generalize across different enforcement regimes. Details on model specification, feature importance, and performance metrics are provided in Appendix C.

To further validate that these exposure measures capture genuine methamphetamine production, I examine whether exposed areas, measured by seizures or machine-learning predictions, show

¹⁹The dropout responses combine the low- θ channel (through π_L), the sector-switching channel, and the high- θ channel (through π_I), so the ratio bounds the population-average density of marginal students rather than isolating it for any one group.

elevated concentrations of chemicals associated with drug synthesis in independent environmental monitoring data.

Investigative reporting has documented that clandestine laboratories routinely discharge toxic waste directly into water sources. A journalistic investigation found that laboratories extract water from rivers and irrigation canals for production, then return contaminated wastewater containing acids, solvents, and heavy metals to the same water bodies (Lab, 2025). Journalists and environmental activists in Sinaloa report finding laboratory equipment—trays, containers, gas cylinders—in rivers, often carried downstream by rainfall from mountain laboratories.

I use water quality measurements from Mexico’s National Water Quality Monitoring Network (RENAMECA), which collects samples at over 5,000 monitoring sites across the country between 2015 and 2019. Methamphetamine production using the P2P (phenyl-2-propanone) method—the dominant synthesis route after Mexico restricted pseudoephedrine imports in 2008—requires mercury as a catalyst and toluene as a solvent. If my exposure measure captures actual drug production, areas with historical laboratory activity should show elevated concentrations of chemical waste in nearby water, particularly mercury and toluene. Appendix F explains the details of this exercise.

Table 9 reports the coefficients from regressions of pollution indicators on different exposure measures. Exposed commuting zones have a higher likelihood of detection of chemical waste, measured by the chemical to biological oxygen demand ratio (COD/BOD), in addition to mercury and toluene. The mercury and toluene regressions control for chemical waste, ensuring that the chemical signature found in the water is specific to methamphetamine production rather than reflecting broad differences in chemical pollution because of differences in industrial activity. Historical seizures of clandestine laboratories are associated with elevated mercury detection rates (+11.9 percentage points) and toluene detection rates (+2.3 percentage points), consistent with residual contamination from P2P synthesis. Importantly, only machine-learning predictions of exposed areas are associated with higher overall chemical waste, and the mercury and toluene effects are more precisely estimated than with the baseline measure based only on historical seizures.

Figure 11 presents estimates of the main results using alternative treatment definitions at two probability thresholds (25% and 50%). The ML-predicted measures classify commuting zones as treated if their predicted probability of production exceeds the threshold. The augmented measures combine the baseline seizure indicator with ML predictions, ensuring 100% recall of observed seizures while adding likely undetected producers.

For middle school dropout, estimates are consistently negative across all specifications. At the 25% threshold, estimates range from -0.018 to -0.025 across ML-predicted and augmented measures, comparable to the baseline estimate of -0.012 . At the 50% threshold, estimates using augmented measures remain similar (-0.018 to -0.022), though ML-predicted measures alone yield smaller and less precise estimates due to the smaller number of treated units (as few as 33 CZs for the Random Forest at 50%).

For high school dropout, results are similarly robust. The baseline estimate of approximately -0.025 is bracketed by ML-based estimates ranging from -0.020 to -0.031 at the 25% threshold. At the 50% threshold, augmented measures continue to show significant negative effects (-0.020 to -0.025), while ML-predicted measures are less precise but directionally consistent.

The stability of results across treatment definitions is reassuring for several reasons. First, it suggests that the baseline findings are not driven by selection into enforcement—areas predicted to have production capacity based on observable characteristics show similar educational responses. Second, the augmented measures, which expand the treated group by 65–70% relative to baseline, yield estimates of similar magnitude, suggesting that undetected producers respond similarly to detected ones. Third, the consistency across different ML algorithms (Random Forest, XGBoost, ensemble) indicates that results do not depend on the specific functional form used to predict production.

9.2 Robustness to Sample Composition

The employment survey used for the main analysis is a rotating panel survey stratified at the state level plus 31 major metropolitan areas, rather than at the commuting zone level. As a result, not all commuting zones are represented in every quarter, yielding an unbalanced panel. Of the 1,226 commuting zones in the analysis sample, only 61 are present across all quarters in the study period. This raises the concern that results could be driven by compositional changes in which commuting zones are sampled over time.

To address this, I re-estimate the main specification restricting the sample to commuting zones present in at least 50% and at least 75% of the analysis quarters. Table 10 presents these results alongside the baseline (full-sample) estimates; the corresponding event studies are shown in Appendix Figure E.1. The restricted samples yield estimates that are similar in magnitude to the baseline, though attenuated and less precise as the panel tightens, suggesting that the findings are not driven by the unbalanced panel structure.

A related limitation of the employment survey analysis is statistical power for heterogeneity tests that require splitting the sample by commuting zone characteristics. With survey-based measures, subgroup analyses—such as examining whether effects differ by baseline cartel presence, distance to the U.S. border, or local industry composition—may be underpowered due to the reduced effective sample size within each cell.

I have recently obtained access to administrative enrollment data from Formato 911, which provides school-level enrollment counts covering the universe of Mexican schools. This administrative data offers three advantages: first, it provides a census rather than a sample, eliminating concerns about survey coverage and panel balance; second, the larger effective sample size will allow for more precise estimates when examining effect heterogeneity across commuting zone characteristics. Third, this will allow me to compare schools closer to clandestine lab sites to provide evidence to

differentiate direct or spillover channels driving increased education investments. Analysis of these data is ongoing and will be incorporated in future versions of this paper.

9.3 Ruling out explanations related to contemporaneous shocks

The post-2019 period coincides with several policy changes and shocks that could potentially confound the estimated effects. I examine three: minimum wage increases, the COVID-19 pandemic, and the rollback of the Prospera conditional cash transfer program.

9.3.1 Minimum Wage Increases

Beginning in 2019, Mexico implemented substantial minimum wage increases, raising from 88.36 pesos per day in 2018 to 102.68 pesos in 2019 (a 16% increase), with further increases reaching 248.93 pesos by 2024. Moreover, municipalities in the Zona Libre de la Frontera Norte (ZLFN)—a strip along the U.S. border—received even larger increases, with the minimum wage doubling to 176.72 pesos in 2019 ([Calderón et al., 2023](#)). Figure E.2 shows the evolution of the minimum wage and a map of the municipalities in ZLFN.

Minimum wage increases could affect educational decisions through multiple channels and could confound my estimates if they differentially affected treated areas. Higher wages raise the opportunity cost of schooling for youth who could enter the labor market, potentially increasing dropout. Alternatively, if minimum wage increases raise household income, they could relax financial constraints and reduce dropout. The net effect depends on which channel dominates and whether treated areas were differentially exposed to these wage changes.

Table 11 examines whether treatment status—history of clandestine labs in the commuting zone—predicts differential exposure to minimum wage increases. I use the share of workers in 2018 whose wages were below the minimum wage that would take effect in 2019 (column 1) or 2021 (column 2) as a measure of potential minimum wage bite.²⁰ Workers in treated commuting zones are less likely to be affected by the minimum wage increases. Exposure to clandestine laboratories results in a 0.4pp decrease in the probability of being affected by the MW increase. These estimates are even smaller when excluding ZLFN municipalities (columns 3–4), with coefficients of -0.1 and -0.3 percentage points respectively. I consider both 2019 and 2021 minimum wage thresholds because the policy involved sustained annual increases throughout my study period: the 2019 measure captures immediate exposure at the start of the post-treatment period, while the 2021 measure captures the broader set of workers whose wages became bound by minimum wage policy as the increases accumulated.

This pattern narrows the channels through which minimum wage policy could explain my results. Since treated areas are *less* affected by minimum wage increases than control areas, confounding

²⁰[Haanwinckel \(2025\)](#) evaluates several approaches to measuring regional minimum wage bite and finds that the fraction affected design is preferable to alternatives such as the effective minimum wage or gap measures, particularly when pre-treatment data is available to test for parallel trends.

would require minimum wage increases to *raise* dropout—through opportunity cost effects dominating income effects—such that control areas experienced larger dropout increases, spuriously generating my finding of relatively lower dropout in treated areas. This channel seems unlikely *a priori* given Mexico’s labor market structure: [Cano-Urbina \(2015\)](#) documents that most young workers’ first labor market experiences occur in the informal sector, where minimum wage laws are not enforced and wages are determined by local market conditions rather than mandated floors. Nevertheless, I test this possibility directly below.

To assess whether minimum wage effects in control areas could explain my results, Table 12 examines the relationship between minimum wage exposure and outcomes *within the control group*. Columns 1–2 show that individuals in CZ more affected by minimum wage increases experienced *reduced* dropout after 2019. Among middle school age youth, a one percentage point increase in the fraction of workers affected by the 2019 minimum wage is associated with a 0.0016 percentage point reduction in dropout, with larger effects using the 2021 threshold (22 percentage points). Effects for high school age youth are similar in magnitude. These results suggest that minimum wage increases reduced dropout in affected areas, consistent with the household income channel dominating the opportunity cost channel. Columns 3–4 confirm this interpretation by showing that minimum wage exposure increased employment among school-age youth. Areas with greater minimum wage bite saw youth employment increase substantially after the policy took effect, suggesting that the income gains from higher wages translated into improved household economic conditions. Columns 5–8 show corresponding increases in household labor income, though estimates are less precise.

These within-control-group results rule out minimum wage increases as an explanation for my main findings. The concern was that minimum wage increases might raise dropout in control areas through opportunity cost effects, spuriously generating my result of relatively lower dropout in treated areas. Instead, I find the opposite: minimum wage exposure *reduced* dropout in control areas. Since control areas were more exposed to minimum wage increases and these increases reduced dropout, any minimum wage effects would work *against* finding lower dropout in treated areas. This implies that my estimates of synthetic drug production effects on dropout are, if anything, attenuated by minimum wage policy rather than driven by it.

Table 13 provides direct robustness tests. First, I examine effects separately for ZLFN municipalities (column 2), which experienced the largest minimum wage increases, versus the rest of the country (column 3). The treatment effect remains negative and significant outside ZLFN (-0.016 for middle school, $p < 0.01$; -0.025 for high school, $p < 0.001$), while effects within ZLFN are smaller and less precisely estimated (-0.006 and -0.017 respectively). This pattern is consistent with minimum wage effects working against my findings in ZLFN, where minimum wage exposure is concentrated.

Second, I directly control for household labor income in column 4. The within-control-group analysis suggests that minimum wage effects on dropout operate primarily through the household income

channel. If differential minimum wage exposure were biasing my estimates, controlling for household income should absorb this variation and change the treatment effect. Instead, the treatment effects remain virtually unchanged when including the income control, confirming that the synthetic drug production effect operates through a distinct channel unrelated to household income differences across treated and control areas.

In summary, minimum wage policy cannot explain my findings. While treated areas have slightly lower shares of workers affected by minimum wage increases, these differences are small in magnitude. More importantly, the within-control-group analysis demonstrates that minimum wage exposure *reduced* dropout, meaning any differential minimum wage effects would work against—not in favor of—finding lower dropout in treated areas. The stability of treatment effects when controlling for household income, combined with the pattern of effects across ZLFN and non-ZLFN areas, further confirms that the synthetic drug production effect operates through channels distinct from minimum wage policy.

9.3.2 COVID-19 pandemic

The COVID-19 pandemic represents an important contemporaneous shock that could potentially confound my estimates. The Chinese fentanyl ban—my identifying variation—occurred in April 2019, just eleven months before Mexico implemented its first pandemic-related restrictions in March 2020. Given this proximity, I must rule out that my results are driven by differential pandemic impacts across treated and control areas rather than by changes in synthetic drug production.

Several features of Mexico’s pandemic response mitigate concerns about differential COVID-19 effects on educational outcomes. Most importantly, school closures were mandated and implemented at the national level by the Ministry of Public Education (SEP), ensuring uniform treatment across all municipalities regardless of local COVID-19 severity or pre-existing synthetic drug production. Furthermore, the SEP implemented policies that effectively "froze" enrollment during the pandemic years: students were not allowed to fail the academic year and were automatically re-enrolled in the same institution until in-person instruction resumed in September 2021.

Nevertheless, I provide three additional pieces of evidence to further support that these results are not driven by the pandemic.

Pre-Pandemic Effects on Time Use

First, I examine an alternative outcome that was not subject to the enrollment freeze: time spent studying. Figure 8 presents event study estimates of exposure to synthetic drug production on weekly minutes spent studying or in school and Table 3 shows the corresponding difference-in-differences estimates. The results show that treated areas experience increases in study time beginning shortly after April 2019—crucially, before the COVID-19 school closures in March 2020. This pre-pandemic response provides direct evidence that behavioral changes in educational invest-

ment preceded the pandemic, consistent with my proposed mechanism rather than pandemic-driven effects.

Balance on Pandemic Severity and Policy Response

Second, I examine whether exposure to synthetic drug production predicts differential pandemic severity or policy responses. This matters for two reasons. Pandemic severity could directly affect educational outcomes through several channels: COVID-19 deaths and hospitalizations may have removed household members who provided childcare or financial support, illness among family members could have increased domestic responsibilities for school-age youth, and economic disruptions from lost workdays or medical expenses may have pushed households to withdraw children from school to contribute income. Additionally, local social distancing policies—even those not directly targeting schools—could have affected dropout decisions by altering labor market conditions. Business closures and mobility restrictions likely reduced employment in formal sectors while potentially increasing the relative returns to informal work, including in illegal economies. If treated areas experienced more severe economic disruptions, young people in these areas may have faced different opportunity costs of remaining in school.

Figure 12 presents standardized differences between areas with and without pre-2015 clandestine lab presence across pandemic-related variables.²¹ The results show no significant differences in COVID-19 severity during 2020: tests, cases, and deaths per capita are statistically indistinguishable between treated and control areas. While I find some differences in the composition of local social distancing policies—with treated areas less likely to implement certain restrictions such as alcohol sales bans, business closures, and mobility checkpoints—these differences do not extend to overall policy intensity. The number of policies implemented is balanced across groups, suggesting similar aggregate economic disruption from pandemic restrictions. Critically, these local policies did not include school closures, which remained under federal jurisdiction and were implemented uniformly nationwide.

Heterogeneity by pandemic intensity

Third, I directly examine whether pandemic severity moderates my treatment effects. Table 14 presents estimates from specifications that interact treatment status with COVID-19 intensity measures and from samples that exclude the hardest-hit areas.

Columns (1)-(2) include interactions between treatment and COVID-19 cases and deaths per capita, respectively. The main treatment effects remain stable in magnitude and significance across both panels. The interaction terms are positive, suggesting that the treatment effect is slightly attenuated in areas with higher COVID-19 intensity—consistent with the pandemic disrupting the mechanisms through which synthetic drug production affects schooling decisions. However, these interactions are

²¹COVID-19 severity comes from publicly available administrative data from Mexico’s Health Ministry (de Salud, 2025) and data on municipal social distancing policies is from Enríquez et al. (2025).

small in magnitude relative to the baseline effects: at the mean of COVID-19 cases, the interaction offsets less than 5% of the main treatment effect for high school dropout. Columns (3)-(6) exclude commuting zones with the highest COVID-19 intensity. The treatment effects persist—and are often larger in magnitude—in these restricted samples, demonstrating that results are not driven by areas with extreme pandemic exposure. For middle school dropout, the effect loses statistical significance when excluding the top 5% and 10% of commuting zones by COVID-19 deaths (columns 5-6), but the point estimates remain similar in magnitude to the baseline, suggesting this is due to reduced statistical power from the smaller sample rather than a substantive change in the relationship.

Taken together, these results indicate that while the COVID-19 pandemic overlapped temporally with my study period, it does not explain the estimated effects of synthetic drug production on educational outcomes. The pre-pandemic emergence of effects on study time, the balance on pandemic severity across treatment groups, and the stability of results to controls for COVID-19 intensity all support interpreting my findings as responses to changes in the synthetic drug economy rather than pandemic-induced disruptions.

9.3.3 Rollback of Prospera

A potential confound is that the government abruptly ended Mexico’s pioneering conditional cash transfer program Prospera in early 2019, just two months before China’s fentanyl ban. Prospera—originally named Progresa when launched in 1997—was among the world’s first conditional cash transfer programs, providing monetary transfers to poor households conditional on children’s school enrollment and attendance. At its peak, the program supported approximately 7 million households, nearly one-fourth of the Mexican population, increasing beneficiaries’ incomes by about 30 percent (Parker and Todd, 2017).

When President López Obrador took office in December 2018, he announced the program’s cancellation in February 2019, replacing it with a new grant program called Becas Benito Juárez (BBJ). The transition was abrupt and disruptive: many beneficiary families reported receiving no payments during the first half of 2019, and the geographic distribution of transfers changed substantially. While total spending under BBJ approached previous Prospera levels by year’s end, progressivity worsened dramatically—households in the poorest localities received less than half the transfers they had received under Prospera, while households in the least marginalized localities saw transfers more than double (Marquez-Padilla et al., 2025).

Marquez-Padilla et al. (2025) estimate the effects of this rollback using a difference-in-differences design comparing localities with high and low pre-rollback Prospera penetration. They find that rollback immediately reduced school enrollment, with effects concentrated among boys of high school age (15–17): enrollment fell by 12 percentage points in fully-penetrated localities, a 17% decline relative to the mean. Effects at middle school ages (12–14) were smaller and not robust to alternative specifications. Rising employment mirrored falling enrollment among high-school-aged boys, suggesting that the school-work tradeoff—particularly in informal construction jobs—drove the dropout response. The authors interpret these findings as consistent with Prospera’s original rationale of relieving liquidity constraints that prevent poor households from financing their children’s education.

Since Prospera targeted poor, rural areas while synthetic drug production infrastructure is more common closer to urban areas with industrial capacity and proximity to the U.S. border and international ports, control commuting zones in my sample have systematically higher Prospera penetration than treated zones. If rollback-induced dropout increases in high-Prospera control areas are not accounted for, this could generate spurious findings of relative dropout decreases in treated areas.

Table 15 directly addresses this concern. Columns 1–2 interact the post-2019 indicator with measures of Prospera penetration, allowing rollback effects to vary with program intensity. I define high-Prospera CZs as those in the top 25th or 50th percentile of the 2017 beneficiary share distribution (the ratio of enrolled households to total households from the 2010 census).

Consistent with [Marquez-Padilla et al. \(2025\)](#), I find that high-Prospera areas experienced significant dropout increases after 2019. CZs in the top quartile of Prospera penetration saw dropout rates rise by 4.5 percentage points for middle school and 7.3 percentage points for high school, relative to low-Prospera areas. Results are similar using the top 50% threshold.

Crucially, the main effect of synthetic drug exposure remains negative after controlling for Prospera-induced variation. In columns 1–2, the coefficient on $\text{Exposure} \times \text{Post}$ —which now captures the effect in low-Prospera CZs—shows that synthetic drug exposure continues to reduce dropout by 0.5–1.2 percentage points for middle school and 1.4–1.9 percentage points for high school. The attenuation relative to baseline estimates is modest, and effects remain statistically significant for most estimations.

The triple interaction ($\text{Exposure} \times \text{Post} \times \text{Prospera}$) reveals a striking pattern: treated CZs in high-Prospera areas experienced *larger* dropout reductions than treated CZs in low-Prospera areas. For middle school, the triple interaction implies that synthetic drug exposure reduced dropout by an additional 6–13 percentage points in high-Prospera CZs relative to low-Prospera CZs. For high school, the additional reduction is 8–24 percentage points. The total effect of synthetic drug exposure in high-Prospera CZs—the sum of $\text{Exposure} \times \text{Post}$ and the triple interaction—is substantially larger than in low-Prospera CZs. Using the top 25% threshold, the total effect in high-Prospera areas represents a 14 percentage point reduction in middle school dropout and a 26 percentage point reduction in high school dropout—effects roughly three times larger than the Prospera rollback effects documented by [Marquez-Padilla et al. \(2025\)](#) and an order of magnitude larger than my baseline estimates.

This pattern is the opposite of what one would expect if Prospera rollback were confounding the results. If the estimated effect of synthetic drug exposure simply reflected dropout increases in high-Prospera control areas, the triple interaction should be zero—not strongly negative. Instead, the results suggest that synthetic drug exposure may partially offset the negative educational consequences of losing transfer income, consistent with skill-intensive criminal opportunities providing alternative incentives for continued schooling.

Columns 3–4 provide a complementary test by excluding the most rural commuting zones, where Prospera penetration was highest and rollback effects most severe. Results are robust to this sample restriction: the effect of exposure to drug production on dropout is consistent with the baseline estimates. This stability confirms that the main findings are not driven by comparisons between treated urban/industrial CZs and control rural CZs differentially affected by Prospera rollback.

Finally, the timing of effects provides further evidence against a Prospera-driven explanation. [Marquez-Padilla et al. \(2025\)](#) document that rollback effects materialized immediately: enrollment gaps between high- and low-Prospera localities emerged abruptly in fall 2019, the first school-year transition after rollback, with no anticipation effects in the preceding months. This timing reflects the nature of the shock—families losing transfer income face immediate pressure to send children

to work rather than school, and dropout decisions concentrate at school-year transitions. In contrast, my event studies (Figure 5) show that effects of synthetic drug exposure emerge gradually, strengthening over the next years rather than appearing as a sharp break in fall 2019. This pattern is consistent with my proposed mechanism: younger cohorts adjusting their educational investments in response to changing skill demands in criminal labor markets, a process that unfolds over multiple school-year decisions rather than as an immediate response to income loss. The divergent dynamics—immediate for Prospera rollback, gradual for synthetic drug exposure—suggest distinct underlying mechanisms. Additionally, my analysis of tertiary enrollment in chemistry-specific STEM fields shows patterns consistent with skill-biased technical change rather than a generic income shock.

10 Conclusion

This paper shows that a skill-biased shift in the technology of an illegal market can affect human-capital decisions of the youth exposed to it. China’s 2019 fentanyl ban made domestic synthesis from uncontrolled pre-precursors the profitable way to produce fentanyl, turning clandestine drug production in Mexico into a chemistry-intensive activity and raising the local demand for chemistry skill in the commuting zones with prior synthetic-drug infrastructure. Comparing those exposed zones against corresponding controls in a difference-in-differences design, I find that exposure reduced school dropout, raised educational attainment, and shifted tertiary enrollment specifically toward chemistry-intensive fields. A model of a local market for chemistry skill organizes these responses and connects them to observable labor-market evidence: the legal wage premium for chemistry-relevant degrees rose and employment in chemistry occupations expanded. At the same time, the observed chemistry workforce became less credentialed and the observable local stock of degree-holding chemists shrank, consistent with the most skilled being drawn into concealed production.

The central contribution is that criminal activity do not always discourage human capital accumulation; the skill content of the criminal task matters. The expansion of this criminal activity increased the demand for educated individuals and, crucially, occurred under low enforcement and stable cartel control, so it arrived without a violence increase. The identified mechanism is the increased demand for skill from the illegal sector, which is transmitted to the legal one through a shared labor market and input linkages between illegal and legal operations. Also, the evidence confirms that the illegal sector is competing for credentialed workers, so that part of the skill upgrading is captured by the criminal organizations themselves.


These findings carry a new lesson for drug policy. Supply-side interdiction is usually evaluated through its effect on prices, quantities, and—increasingly—violence. This evidence points to a further margin: interdiction can push criminal organizations *up the skill ladder*. Before the ban, fentanyl was for the most part synthesized in Chinese chemical firms—a legal, industrial process—

and only smuggled in finished form in small shipments; the chemical skill was outside the criminal organizations entirely. Banning the finished drug in China did not eliminate its demand in the U.S., instead it relocated synthesis into clandestine Mexican laboratories, moved it back to precursors and pre-precursors, and forced the criminal organizations to internalize the chemical expertise—and the accident and exposure risks—that were absorbed by legal firms before. Interdiction can therefore leave criminal organizations more technologically sophisticated than it found them, and potentially make these organizations harder to dislodge in the long run.

The welfare consequences are look different at the individual and the aggregate level. Privately, illegal-bound youth who were always going to enter illicit production likely benefit from being pushed into getting more education, in so far as this means going into a less dangerous position with a higher pay like switching from working as a soldier into being a lab cook. The young people induced to switch from legal to illegal work take on arrest risk and forgo a legal career, so their private gains are far more ambiguous and plausibly negative once the risk and the criminalization are counted. For the much larger group who simply acquire more schooling in response to the higher legal return, the effect is likely positive—provided the local labor market can absorb the additional chemistry skills rather than leaving them idle or mismatched.

At the aggregate level, more education and the human-capital gains documented here are bound up with the strengthening of the criminal organizations that generated the demand: the same process that keeps youth in school also supplies the illicit sector with a more skilled workforce and a more advanced production capability. Whether the net effect is positive will depend on how one weighs a durable increase in local human capital against a more entrenched and more sophisticated criminal economy, and whether the education gains are durable in this context.

Figure 1: Illegal drug seizures in Mexico




Notes: Panels (a) and (b) show the number of clandestine laboratory and kilogram amount of fentanyl seizures by year, respectively. Vertical shaded areas indicate Mexican presidential administrations: VFQ (Vicente Fox Quesada, 2000-2006), FCH (Felipe Calderón Hinojosa, 2006-2012), EPN (Enrique Peña Nieto, 2012-2018), and AMLO (Andrés Manuel López Obrador, 2018-2024). The vertical line marks 2019, when synthetic drug production became widespread. Panel (c) shows the spatial distribution of clandestine laboratory seizures from 2000-2015 at the municipality level.

Source: Author's calculations based on Mexican Secretariat of National Defense (SEDENA) seizure records obtained via Insight Crime.

[↔ Return to text](#)

Figure 2: Evolution of overdoses in the United States by type of drug




Notes: The figure shows overdose death rates per 100,000 people by drug type from 2002-2023. Synthetic opioids include fentanyl and its analogs. Natural opioids include morphine, codeine, and semi-synthetic opioids such as oxycodone and hydrocodone. The first shaded region indicates the introduction of fentanyl into U.S. supplies. The second shaded region marks China’s fentanyl ban in 2019, which coincided with accelerated growth in synthetic opioid deaths.

Source: National Institute on Drug Abuse (NIDA), based on CDC Wonder, Multiple Cause of Death data.

[↔ Return to text](#)


Figure 3: Spatial distribution of treatment and post-treatment drug production



Notes: The map shows the geographic distribution of the main treatment variable and subsequent drug production activity. Gray shading indicates treatment status: commuting zones with high exposure (dark gray) had clandestine laboratory seizures before 2015, while low exposure zones (light gray) had no pre-2015 labs. Purple circles show fentanyl seizures occurring after the Chinese fentanyl ban, with circle size proportional to kilograms seized. Green diamonds indicate post-2019 clandestine laboratory seizures. Correlations demonstrate that drug production activity persists in treated regions: areas with pre-2015 exposure experience higher rates of both fentanyl seizures and new laboratory discoveries. *Source:* Author's calculations based on Mexican Secretariat of National Defense (SEDENA) press releases and seizure records.

[↔ Return to text](#)

Figure 4: Balance on covariates across treatment and control CZ




Notes: Each point reports the mean of the indicated baseline (pre-2019) commuting-zone characteristic, standardized to mean zero and unit standard deviation, shown separately for treated commuting zones (any methamphetamine laboratory seized before 2015) and control zones; horizontal bars are 95% confidence intervals. Asterisks denote characteristics for which the treated-control difference is statistically significant in a two-sample *t*-test. Characteristics are grouped into geography and access, illegal activities, demographics, education, and employment and firms.

[↔ Return to text](#)

Figure 5: Event Study of Exposure to Synthetic Drug Production on School Dropout


All



Notes: The figure plots event-study coefficients from a difference-in-differences regression of an indicator for not being enrolled in school on commuting-zone exposure interacted with period indicators, omitting the period immediately before the fentanyl ban as the reference. Panel (a) is middle school age (12–14) and panel (b) high school age (15–17). Exposure is an indicator for any methamphetamine laboratory seized in the commuting zone before 2015. All specifications include commuting zone, year-quarter, and quarter fixed effects and control for gender, age, and household education; observations are weighted by survey weights. Green (pink) markers denote pre- (post-) treatment periods, and bands are 95% confidence intervals based on standard errors clustered at the commuting zone level. The shaded region marks the COVID-19 school-closure period.
Source: Author's calculations based on ENOE.

[↔ Return to text](#)


Figure 6: Event studies of the Synthetic Drug Expansion on Primary Age School Dropout



Notes: Placebo event study for primary school age children (6–11), who are not at the margin of leaving school for work. The dependent variable is an indicator for not being enrolled in school. The specification matches Figure 5: commuting zone, year-quarter, and quarter fixed effects, controls for gender, age, and household education, survey weights, and the period before the fentanyl ban as the reference. Bands are 95% confidence intervals from standard errors clustered at the commuting zone level.
Source: Author's calculations based on ENOE.

[↔ Return to text](#)

Figure 7: Impact of synthetic drug expansion on school dropout, results by quarter




Notes: Event-study estimates of the effect of exposure on school dropout, estimated separately for each calendar quarter of the survey. The dependent variable is an indicator for not being enrolled in school; rows show middle school age (12–14) and high school age (15–17), and columns show Q4 (Oct–Dec, start of the Mexican academic year), Q1 (Jan–Mar), and Q2 (Apr–Jun, end of the year). Exposure is an indicator for any methamphetamine laboratory seized in the commuting zone before 2015, interacted with year indicators relative to the last pre-ban observation of each quarter. All specifications include commuting zone and year fixed effects and control for gender, age, and household education; observations are weighted by survey weights. Bands are 95% confidence intervals from standard errors clustered at the commuting zone level; the shaded region marks the COVID-19 school-closure window.

Source: Author's calculations based on ENOE.

[↔ Return to text](#)

Figure 8: Event study of exposure to synthetic drug production on weekly time spent studying




Notes: Event-study estimates of the effect of exposure on weekly time spent studying among school-age youth (12–17), from a Poisson (PPML) regression; coefficients are proportional effects (log points). Exposure is an indicator for any methamphetamine laboratory seized in the commuting zone before 2015, interacted with period indicators relative to the period before the fentanyl ban. The model includes commuting zone, year-quarter, and quarter fixed effects and controls for gender, age, and household education; observations are weighted by survey weights. Bands are 95% confidence intervals from standard errors clustered at the commuting zone level. The shaded region marks the COVID-19 school-closure period.

Source: Author's calculations based on the ENOE time-use module.

[↔ Return to text](#)

Figure 9: Effect of Exposure to Synthetic Drug Production on Reaching the Final Grade of High School




Notes: The figure plots coefficients from a cohort event study estimated on the Formato 911 school census. The dependent variable is the share of a middle-school entry cohort in a commuting zone that remains continuously enrolled through the final grade of high school (a six-grade trajectory). Cohorts are indexed by the grade they were enrolled in when the fentanyl ban was announced in April 2019 (x -axis; age $\approx$ grade + 5), and the specification includes commuting-zone and grade-at-ban fixed effects, weights cohorts by size, and clusters standard errors at the commuting-zone level. The omitted category is the cohort enrolled in the final grade at the ban (the last to reach it before the shock). Green markers denote cohorts that reached the final grade before the ban and pink markers those that reached it in 2019 or later; the vertical dashed line marks the ban, and the shaded region marks cohorts that reached the final grade during or after the COVID-19 school closures. The single unshaded post-ban cohort (grade 11 at the ban) reached the final grade in September 2019—after the ban but before the closures. Bands are 95% confidence intervals and dashed lines 90% confidence intervals.

Source: Author's calculations based on Formato 911.

[↔ Return to text](#)

Figure 10: Impact of the Synthetic Drug Expansion on Tertiary enrollment




Notes: Event-study estimates of the effect of exposure on tertiary (university) enrollment by field of study, at the commuting-zone level. Panels show total enrollment, chemistry-intensive STEM+H programs, and STEM+H programs without chemistry. Estimates are from a Poisson (PPML) regression of field-specific enrollment counts on commuting-zone exposure interacted with year indicators (reference: 2018, the year before the ban), with commuting zone and year fixed effects; coefficients are proportional effects (log points). Exposure is an indicator for any methamphetamine laboratory seized in the commuting zone before 2015. Bands are 95% confidence intervals from standard errors clustered at the commuting zone level.

Source: Author's calculations based on ANUIES tertiary enrollment data.

[↔ Return to text](#)


Figure 11: DiD estimates of alternative treatment measures



Notes: Difference-in-differences estimates of the effect of exposure on school dropout (Panel A) and tertiary enrollment (Panel B) under alternative definitions of treatment. Within each panel, the x-axis indicates the exposure measure: the baseline seizure-based indicator (any methamphetamine laboratory seized before 2015), and machine-learning predictions of latent production from a random forest (RF), gradient boosting (XGB), and their ensemble. Point shape denotes the treatment type—seizure-based, ML-predicted (a commuting zone is treated if its predicted probability exceeds the threshold), or augmented (the maximum of the seizure-based and ML-predicted indicators); color denotes the classification threshold (.2, .3, .4, .5). Panel A reports linear-probability DiD coefficients for middle and high school dropout (commuting zone, year-quarter, and quarter fixed effects; controls for gender, age, and household education). Panel B reports Poisson (PPML) coefficients—proportional effects—for total enrollment, chemistry-intensive STEM+H, and STEM+H without chemistry (commuting zone and year fixed effects). Bars are 95% confidence intervals from standard errors clustered at the commuting zone level.

[↔ Return to text](#)

Figure 12: Balance on COVID-19 severity and local policy response by exposure to synthetic drug production




Notes: The figure assesses balance in COVID-19 severity and the local policy response across commuting zones by treatment status. Each point is the mean of the indicated measure (e.g., 2020 COVID-19 cases and deaths per capita and local policy indicators), standardized to mean zero and unit standard deviation, shown separately for treated commuting zones (any methamphetamine laboratory seized before 2015) and control zones; horizontal bars are 95% confidence intervals. Asterisks denote treated-control differences significant in a two-sample t -test.

Source: Author's calculations.

[↔ Return to text](#)


Figure 13: Exposure and the intentional-homicide rate



Notes: Event-study estimates of exposure on intentional homicides at the commuting-zone $\times$ year level, 2015–2024, relative to the year before the ban ($t = -1$). Estimated by Poisson pseudo-maximum-likelihood with a log-population offset, so coefficients are proportional changes in the per-capita homicide rate; population is the annual commuting-zone total from CONAPO municipal projections. Specification includes commuting-zone and year fixed effects; standard errors clustered by commuting zone. Robustness to alternative population measures and to the high-violence subsample is reported in Appendix Table G.1 and Appendix Figure G.1.

[↔ Return to text](#)


Figure 14: Exposure and survey non-response: household refusal rate



Notes: Event-study estimates of exposure on the household refusal rate—the share of inhabited dwellings whose occupants decline the ENOE interview—at the commuting-zone $\times$ year level, relative to the year before the ban ($t = -1$). Because ENOE expansion factors are zero for non-respondents, the panel is unweighted. Specification includes commuting-zone and year fixed effects, a seasonal (quarter) fixed effect, and commuting-zone linear trends; standard errors clustered by commuting zone. Green markers are pre-treatment estimates, purple post-treatment; shaded bands are 95% confidence intervals. The shaded vertical region marks the COVID-19 school-closure window.

[↔ Return to text](#)


Figure 15: Event studies: the price and quantity of chemistry skill



Notes: Event studies of the two headline labor-market outcomes, relative to the year before the ban ($t = -1$). Panel (a) is the change in the wage premium on a STEM degree, from a triple-difference on log hourly wages; the chemistry-specific estimate is in Table 6, and the broader STEM series is shown here as the cleaner parallel-trends check. Panel (b) is the share of employment in chemistry occupations, estimated by Poisson pseudo-maximum-likelihood at the commuting-zone $\times$ year level with a log-employment offset. Both specifications include commuting-zone and year (survey-wave) fixed effects; standard errors clustered by commuting zone. Green markers are pre-treatment estimates, purple post-treatment; shaded bands are 95% confidence intervals; the shaded vertical region marks the COVID-19 school-closure window.

[↔ Return to text](#)

Figure 16: Exposure and the STEM-degree share of the fixed 2019 cohort



Notes: Event study of the share of the fixed 2019 cohort (aged 30–54 in 2019) holding a STEM degree, relative to the year before the ban ($t = -1$). Poisson pseudo-maximum-likelihood at the commuting-zone $\times$ year level with an offset for log cohort size, commuting-zone and year fixed effects; standard errors clustered by commuting zone. Green markers are pre-treatment estimates, purple post-treatment; shaded bands are 95% confidence intervals; the shaded vertical region marks the COVID-19 school-closure window.

[↔ Return to text](#)

Table 1: Effect of Exposure to Synthetic Drug Production on School Dropout

	Middle school				High school			
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
Exposure $\times$ Post	-0.017** (0.005)	-0.017** (0.006)	-0.016** (0.005)	-0.011+ (0.006)	-0.021* (0.009)	-0.022* (0.009)	-0.024** (0.008)	-0.027** (0.009)
Exposure $\times$ Post $\times$ Female				-0.008 (0.006)				0.006 (0.011)
Female		-0.009*** (0.001)	-0.009*** (0.001)	-0.007*** (0.002)		-0.028*** (0.002)	-0.029*** (0.002)	-0.017*** (0.003)
Age		0.021*** (0.002)	0.020*** (0.002)	0.020*** (0.002)		0.054*** (0.003)	0.053*** (0.003)	0.053*** (0.003)
Household avg. education			-0.003*** (0.000)	-0.003*** (0.000)			-0.010*** (0.001)	-0.010*** (0.001)
N	665,100	665,100	665,100	665,100	674,437	674,437	674,437	674,437
Mean of dep. var.	0.059	0.059	0.059	0.059	0.245	0.245	0.245	0.245

Notes: The dependent variable is an indicator for not being enrolled in school among individuals of the relevant school age. Exposure is an indicator for any methamphetamine laboratory seized in the commuting zone before 2015. Post is an indicator for periods after April 2019. Household avg. education is the average years of schooling of adults older than 25 in the individual's household. All specifications include commuting zone, year-quarter, and quarter fixed effects. Standard errors clustered at the commuting zone level are reported in parentheses. + $p < 0.1$, * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

[↩ Return to text](#)

Table 2: Impact of Exposure to Synthetic Drug Production on School Dropout by quarter

	Middle school			High school		
	Oct-Dec	Jan-Mar	Apr-Jun	Oct-Dec	Jan-Mar	Apr-Jun
Exposure $\times$ Post	-0.016*** (0.005)	-0.018** (0.006)	-0.012* (0.006)	-0.014 (0.009)	-0.024** (0.009)	-0.033*** (0.007)
Female	-0.007*** (0.001)	-0.008*** (0.002)	-0.012*** (0.002)	-0.024*** (0.003)	-0.029*** (0.003)	-0.034*** (0.003)
Age	0.019*** (0.002)	0.021*** (0.002)	0.019*** (0.002)	0.055*** (0.003)	0.053*** (0.003)	0.050*** (0.002)
HH schooling (25+)	-0.003*** (0.000)	-0.003*** (0.000)	-0.003*** (0.000)	-0.009*** (0.001)	-0.010*** (0.001)	-0.010*** (0.001)
N	214,645	230,717	219,738	218,077	234,183	222,177
Mean of dep. var.	0.054	0.058	0.066	0.230	0.244	0.263

Notes: This table reports difference-in-differences estimates of the effect of increased synthetic drug production on school dropout, estimated separately by quarter. The dependent variable is an indicator for not being enrolled in school among individuals of the relevant school age. Middle school age is 12–14 years; high school age is 15–17 years. Q4 (October–December) corresponds to the first quarter of the Mexican academic year, which runs August–July; Q2 (April–June) corresponds to the end of the academic year. Exposure is an indicator for any methamphetamine laboratory seized in the commuting zone before 2015. Post is an indicator for periods after April 2019. HH schooling (25+) is the average years of schooling of adults older than 25 in the individual’s household. All specifications include commuting zone and year fixed effects and control for gender, age, and household education. Standard errors clustered at the commuting zone level are reported in parentheses. + $p < 0.1$, * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

[↔ Return to text](#)

Table 3: Impact of Exposure to Synthetic Drug Production on Time Use

Statistic	Study (Any)	Study (Min)	Care (Min)	Shopping (Min)	Chores (Min)	Volunteer (Min)
<i>Middle School</i>						
Exposure $\times$ Post	0.0211*** (0.0067)	0.0226** (0.0104)	0.0558 (0.1302)	-0.0385 (0.0467)	-0.0493*** (0.0181)	-0.1661 (0.2308)
Model	LPM	Poisson	Poisson	Poisson	Poisson	Poisson
Control Mean	0.8579	1793.3248	45.6872	29.1952	358.6615	1.2775
N	655,771	654,047	651,645	654,418	653,028	573,318
<i>High School</i>						
Exposure $\times$ Post	0.0325*** (0.0076)	0.0504*** (0.0134)	-0.0218 (0.0663)	-0.0334 (0.0472)	-0.0641*** (0.0149)	0.0670 (0.2293)
Model	LPM	Poisson	Poisson	Poisson	Poisson	Poisson
Control Mean	0.6623	1445.3910	93.3901	39.6596	472.8108	2.5293
N	674,457	672,107	670,775	673,212	670,881	614,398

Notes: Difference-in-differences estimates of the effect of exposure to synthetic drug production on weekly time use among school-age youth (12–17 years). Study (Any) is an indicator for reporting any time spent studying and is estimated by a linear probability model; the remaining (minutes) outcomes are estimated by Poisson, so the reported coefficients are proportional effects (semi-elasticities). The control mean is reported in minutes for reference. Exposure indicates any methamphetamine laboratory seized in the commuting zone before 2015; Post indicates years after 2019. All models include commuting zone, year-quarter, and quarter fixed effects and control for gender, age, and household education; observations are weighted by survey weights. Standard errors clustered at the commuting zone level in parentheses. * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

[↔ Return to text](#)

Table 4: Impact of the Synthetic Drug Expansion on tertiary enrollment

	Enrollment			
	STEM+H			Non-STEM+H
	All fields	Chemistry	Non-Chemistry	
<i>Panel A. All</i>				
Exposure $\times$ Post	0.000 (0.022)	0.057* (0.023)	-0.016 (0.020)	-0.028 (0.026)
N	9,304	8,944	8,600	9,280
Control mean	2518.300	744.700	404.400	1435.000
<i>Panel B. Female</i>				
Exposure $\times$ Post	-0.010 (0.021)	0.061** (0.022)	-0.049+ (0.027)	-0.030 (0.026)
N	9,304	8,944	8,600	9,280
Control mean	1291.700	336.100	184.400	801.500
<i>Panel C. Male</i>				
Exposure $\times$ Post	0.010 (0.024)	0.048+ (0.026)	0.010 (0.019)	-0.024 (0.027)
N	9,304	8,944	8,600	9,280
Control mean	1226.600	408.600	220.000	633.500

Notes: All models include commuting zone and year fixed effects. Standard errors clustered at the commuting zone level. + $p < 0.1$, * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

[↔ Return to text](#)

Table 6: Effects of synthetic drug production on returns to education and employment

	(1)	(2)	(3)	(4)	(5)	(6)
<i>Panel A. Effects on wages and returns to education degrees</i>						
	Dependent variable: Hourly wage					
		Degree group interaction				
					STEM by chemistry content	
	Overall wage	Any tertiary	Non-STEM	STEM	No	Yes
Exposure $\times$ Post	-0.0121 (0.0143)	-0.0211 (0.0206)	-0.0195 (0.0187)	-0.0135 (0.0152)	-0.0129 (0.0147)	-0.0127 (0.0148)
Exposure $\times$ Post $\times$ Degree _g		0.0407** (0.0197)	0.0318 (0.0195)	0.0501** (0.0222)	0.0234 (0.0229)	0.0589** (0.0255)
Observations (ind. $\times$ year)	3,655,535	3,655,535	3,655,535	3,655,535	3,655,535	3,655,535
Mean (control)	3.388	3.388	3.388	3.388	3.388	3.388
Baseline premium (log pts)		0.733	0.642	0.514	0.468	0.520
No. of commuting zones	1,230	1,230	1,230	1,230	1,230	1,230
<i>Panel B. Effects on employment and share of employment by occupation</i>						
	Dependent variable: employment (PPML)					
		By occupation group				
					STEM by chemistry content	
	All	Professional	Non-STEM	STEM	No	Yes
Exposure $\times$ Post	-0.0015 (0.0034)	0.0130 (0.0096)	0.0041 (0.0107)	0.0733*** (0.0283)	0.0514 (0.0326)	0.1762*** (0.0606)
Observations (CZ $\times$ year)	8,955	8,768	8,756	6,771	6,484	4,434
Mean (control)	0.7568	0.1931	0.1720	0.0215	0.0179	0.0036
No. of commuting zones	1,230	1,152	1,149	729	682	435

Notes: Difference-in-differences estimates; standard errors clustered by commuting zone in parentheses. *Exposure* equals one for commuting zones with a methamphetamine laboratory seized before 2015; *Post* equals one from 2019 (China's fentanyl ban). Sample: prime-age (25–54). *Panel A* reports OLS estimates on the log hourly wage: column (1) is the overall wage effect, and columns (2)–(6) add a degree-group (*g*) interaction whose triple interaction (*Exposure* $\times$ *Post* $\times$ *Degree_g*) is the change in that degree's wage premium in treated commuting zones after the shock; all columns control for sex and a quadratic in age, include commuting-zone and survey-wave (year $\times$ quarter) fixed effects, and are weighted by survey weights. *Panel B* reports Poisson (PPML) estimates at the commuting-zone $\times$ year level: column (1) is the employment rate (log-population offset) and columns (2)–(6) are the share of employment in each occupation group (log-employment offset), so coefficients are proportional changes; all include commuting-zone and year fixed effects. The number of commuting zones falls across the columns of Panel B because the Poisson estimator drops commuting zones with no within-zone variation in the outcome—a zone with no employment in a given occupation group in any year is perfectly predicted by its fixed effect and contributes nothing to the likelihood. In both panels columns (5)–(6) split STEM by chemistry content; column (6) is chemistry: chemistry-course fields in Panel A (including pharmacy) and chemistry occupations in Panel B (chemists, chemical engineers, and chemical-plant operators).

* $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$.

[↔ Return to text](#)


Table 7: Effects on the composition of the workforce: education and geographic origin

	(1)	(2)	(3)	(4)	(5)	(6)
<i>Panel A. Educational composition of the workforce (average years of schooling)</i>						
	Dependent variable: average years of schooling					
	By occupation group					
	STEM by chemistry content					
	All	Professional	Non-STEM	STEM	No	Yes
Exposure $\times$ Post	-0.0365 (0.0399)	0.0146 (0.0374)	0.0224 (0.0395)	-0.1518*** (0.0586)	-0.0843 (0.0520)	-0.2849** (0.1370)
Observations (CZ $\times$ year)	9,657	8,196	8,138	3,729	3,328	1,868
Mean (control)	10.879	14.814	14.666	15.902	16.217	14.361
No. of commuting zones	1,235	1,162	1,159	740	694	443
<i>Panel B. Geographic origin of the workforce (out-of-state share)</i>						
	Dependent variable: out-of-state (migrant) share					
	By occupation group					
	STEM by chemistry content					
	All	Professional	Non-STEM	STEM	No	Yes
Exposure $\times$ Post	-0.0089* (0.0047)	-0.0000 (0.0057)	-0.0001 (0.0053)	-0.0023 (0.0124)	-0.0023 (0.0131)	-0.0021 (0.0175)
Observations (CZ $\times$ year)	9,657	8,199	8,141	3,730	3,328	1,870
Mean (control)	0.2091	0.2153	0.2100	0.2584	0.2638	0.2321
No. of commuting zones	1,235	1,162	1,159	740	694	443

Notes: Difference-in-differences estimates; standard errors clustered by commuting zone in parentheses. Both panels are commuting-zone $\times$ year weighted OLS (weights = workforce size) with commuting-zone and year fixed effects, estimated on prime-age (25–54) employed workers. *Exposure* equals one for commuting zones with a methamphetamine laboratory seized before 2015; *Post* equals one from 2019 (China’s fentanyl ban). Each outcome is computed over the workforce of the occupation group heading the column: *Panel A* is the average years of schooling of that group’s workers (the coefficient is a change in years), and *Panel B* is the share of that group’s workers born in a state other than their state of residence (a change in percentage points). Columns follow the occupation grouping of Table 6, Panel B: (1) all employed, (2) professional occupations (ISCO-08 major group 2), (3) non-STEM professional, (4) STEM, and (5)–(6) STEM split by chemistry content, with (6) the chemistry occupations (chemists, chemical engineers, and chemical-plant operators). The number of commuting zones falls across columns because thinner occupation groups are present in fewer zones. * $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$.

[↩ Return to text](#)

Figure 17: Model-implied illegal-sector magnitudes from the estimated legal chemistry premium
 $(\widehat{d \ln \pi_L} = 0.0589)$



Notes: Both panels invert the difference-in-differences estimate of the legal chemistry premium $(\widehat{d \ln \pi_L} = 0.0589)$ through the market-clearing conditions of Section 3. Panel (a) plots the implied rise in the *illegal* chemistry premium against the reallocation elasticity κ_L ; Panel (b) plots the implied growth in illegal demand for chemistry skill against the illegal-to-legal chemist ratio r (log scale). The solid line holds the remaining parameters at central values ($\sigma_L = 1$, $\alpha_L = 0.10$, $\lambda = 0.8$, $\sigma_I + \alpha_I = 0.85$); the shaded band traces the legal demand elasticity over $\sigma_L \in [0.75, 1.5]$ and is a parameter-robustness range, not a confidence interval. In Panel (b), the dashed lines mark the growth of detected clandestine-laboratory (Δ lab) and fentanyl (Δ fentanyl) seizures over the study window (average 2020–2023 relative to average 2015–2018, dropping the 2019 transition year), and the solid points their implied r ; because seizures undercount production, these growth factors are lower bounds and the implied r upper bounds.

Source: Author's calculations.

[↔ Return to text](#)

Table 8: Effects on the stock of degree holders, aged 30–54 at the time of the shock

	(1)	(2)	(3)	(4)	(5)
Dependent variable: share of the cohort holding a degree					
	By field of study				
	Any tertiary	Non-STEM	STEM	STEM by chemistry content	
				No	Yes
Exposure $\times$ Post	0.0020 (0.0149)	0.0067 (0.0164)	-0.0366* (0.0193)	-0.0253 (0.0280)	-0.0461** (0.0226)
Observations (CZ $\times$ year)	9,236	9,260	7,998	6,873	7,282
Control share (degree)	0.2243	0.2211	0.0548	0.0207	0.0343
No. of commuting zones	1,093	1,097	838	671	738

Notes: Poisson (PPML) difference-in-differences estimates of the share of a fixed birth cohort holding a degree; standard errors clustered by commuting zone in parentheses. The cohort is people aged 30–54 *in 2019*, followed across 2012–2024 by age in 2019, so their attainment is predetermined and a decline cannot reflect new schooling—only mobility or observability. Estimated on commuting-zone $\times$ year cells with an offset for log cohort size, commuting-zone and year fixed effects; *unconditional* on employment. *Exposure* equals one for commuting zones with a methamphetamine laboratory seized before 2015; *Post* equals one from 2019. Column (1) is any tertiary degree; columns (2)–(5) are by field of study, with (4)–(5) splitting STEM by chemistry content ((5) = chemistry-course fields, including pharmacy). Any tertiary (1) and non-STEM (2) are placebos: a decline concentrated in STEM and chemistry, with the placebos flat, rules out generic educated out-migration or survey non-response, which would fall on all fields alike. * $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$.

[↔ Return to text](#)

Table 9: Validating Exposure Measures: Future Drug Lab Seizures and Methamphetamine Production Signatures in Water

	Lab Seizures Post-2015		COD/BOD (std.)		Mercury Detection		Toluene Detection	
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
Observed Exposure	0.086*** (0.012)	0.096*** (0.013)	-0.131** (0.040)	-0.138*** (0.041)	0.119*** (0.017)	0.104*** (0.016)	0.023*** (0.005)	0.020*** (0.005)
Pred. Exposure RF	0.047*** (0.007)		0.201*** (0.036)		0.050*** (0.012)		0.016*** (0.004)	
Pred. Exposure XGB		0.031*** (0.007)		0.215*** (0.038)		0.075*** (0.012)		0.022*** (0.004)
N sites			-0.002*** (0.000)	-0.002*** (0.000)	-0.000*** (0.000)	-0.000*** (0.000)	-0.000*** (0.000)	-0.000*** (0.000)
COD/BOD (std.)					-0.018*** (0.003)	-0.019*** (0.003)	-0.004** (0.001)	-0.004*** (0.001)
Mean DV	0.015	0.015	0.000	0.000	0.312	0.312	0.073	0.073
N	17700	17700	10111	10111	10111	10111	8326	8326
R ²	0.098	0.091	0.040	0.041	0.225	0.230	0.040	0.041

Notes: Each column reports estimates from an OLS regression at the commuting zone (CZ)–year level for 2015–2024 (columns 1–2) and 2015–2020 (columns 3–8). All specifications include year fixed effects and cluster standard errors at the CZ level. Columns 1–2 estimate the probability of a drug lab seizure in CZ c in year t as a function of both observed and ML-predicted pre-2015 lab exposure. Columns 3–8 use water quality outcomes measured at monitoring stations aggregated to the CZ level using inverse-distance contiguity weights. The COD/BOD ratio (columns 3–4) is standardized to mean zero and unit standard deviation. Mercury and toluene detection rates (columns 5–8) measure the weighted share of monitoring sites where concentrations exceeded detection limits, controlling for the COD/BOD ratio to absorb general industrial pollution. All water quality regressions (columns 3–8) are weighted by site-CZ contiguity weights; columns 1–2 are unweighted. "Observed Exposure" equals one if any synthetic drug laboratory was seized in the CZ before 2015. "Pred. Exposure RF" and "Pred. Exposure XGB" equal one if the random forest or gradient-boosted ensemble, respectively, assigns a predicted probability of lab presence above 0.25. Odd- and even-numbered columns within each outcome panel differ only in the ML algorithm used for predicted exposure. + $p < 0.1$, * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

[↔ Return to text](#)

Table 10: Robustness to commuting zone panel balance

Statistic	Full	Balanced 50	Balanced 75
Middle school			
Exposure × Post	-0.0155*** (0.0051)	-0.0134*** (0.0052)	-0.0094* (0.0050)
Control mean	0.059	0.054	0.047
N	665,100	639,338	607,051
High school			
Exposure × Post	-0.0245*** (0.0076)	-0.0221*** (0.0077)	-0.0151* (0.0078)
Control mean	0.245	0.233	0.218
N	674,437	650,705	619,275

Notes: Difference-in-differences estimates of the effect of exposure on school dropout, re-estimated on progressively more balanced commuting-zone panels. "Full" is the baseline sample; "Balanced 50" ("Balanced 75") restricts to commuting zones present in at least 50% (75%) of survey waves. The dependent variable is an indicator for not being enrolled in school. Exposure is an indicator for any methamphetamine laboratory seized in the commuting zone before 2015. All specifications include commuting zone, year-quarter, and quarter fixed effects and control for gender, age, and household education; observations are weighted by survey weights. Standard errors clustered at the commuting zone level in parentheses. + $p < 0.1$, * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

[↔ Return to text](#)

Table 11: Exposure to synthetic drug production and share of workers affected by minimum wage increases

	Dep. variable: Worker is affected by MW increase			
	Full sample		Excluding ZLFN	
	2019 MW	2021 MW	2019 MW	2021 MW
	(1)	(2)	(3)	(4)
Exposure	-0.004*** (0.000)	-0.008*** (0.001)	-0.001*** (0.000)	-0.003*** (0.001)
Num.Obs.	705297	705297	664183	664183

+ p < 0.1, * p < 0.05, ** p < 0.01, *** p < 0.001

Notes: This table reports OLS estimates of the correlation between exposure to clandestine drug labs and the share of workers affected by minimum wage increases. The dependent variable is a dummy equal to one if the worker's job is formal and their daily wage in 2018 was between the current minimum wage that would take effect in 2019 (columns 1 and 3) or 2021 (columns 2 and 4). Columns 1-2 use the full sample; columns 3-4 exclude municipalities in the Zona Libre de la Frontera Norte (ZLFN), which experienced larger minimum wage increases. Standard errors in parentheses. Observations are weighted using survey expansion factors. + p < 0.1, * p < 0.05, ** p < 0.01, *** p < 0.001.

Data: ENOE 2018.

[↔ Return to text](#)

Table 12: Minimum wage increases and dropout in the control group

term	Dependent variable:							
	School Dropout		Works		HH Labor Income			
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
Panel A: Middle school age								
FA by 2019 MW × Post	-0.164+ (0.099)		-0.047 (0.075)		0.692*** (0.201)		0.830*** (0.168)	
FA by 2021 MW × Post		-0.222*** (0.072)		-0.103* (0.048)		-0.132 (0.281)		0.222 (0.189)
Num.Obs.	394893	394893	394895	394895	322977	322977	310572	310572
Panel B: High school age								
FA by 2019 MW × Post	-0.129 (0.116)		-0.220 (0.269)		0.873*** (0.232)		1.141*** (0.186)	
FA by 2021 MW × Post		-0.264** (0.094)		-0.177+ (0.106)		-0.386 (0.319)		0.011 (0.250)
Num.Obs.	397663	397663	397670	397670	321500	321500	297019	297019

Notes: This table reports difference-in-differences estimates of the effect of exposure to minimum wage increases on school dropout and household labor income. The dependent variable in columns 1-2 is a dummy for dropping out of middle school (Panel A) or high school (Panel B). The dependent variable in columns 3-4 is equal to one if the individual is currently working. The dependent variable in columns 5-8 is the log of household labor income per capita (col. 5-6), and per capita household labor income considering only income of members aged 25+ (cols 7-8). All columns restrict to individuals of middle school age (Panel A) or high school age (Panel B). All specifications include commuting zone, year-quarter, and quarter fixed effects, and control for sex, age, and household schooling of members aged 25+. Observations are weighted using normalized survey weights. Standard errors clustered at the commuting zone level are reported in parentheses. + p < 0.1, * p < 0.05, ** p < 0.01, *** p < 0.001.

[↔ Return to text](#)

Table 13: Robustness to minimum wage increases

term	Dependent variable: School dropout			
	(1)	(2)	(3)	(4)
Panel A: Middle school age				
Exposure $\times$ Post	-0.016*** (0.005)	-0.006 (0.009)	-0.016*** (0.006)	-0.015* (0.006)
HH Income pc (log)				-0.000 (0.001)
Num.Obs.	665100	42167	622933	494326
Sample	Full	ZLFN	Not ZLFN	Full
Panel B: High school age				
Exposure $\times$ Post	-0.024*** (0.008)	-0.017+ (0.009)	-0.025*** (0.008)	-0.023* (0.011)
HH Income pc (log)				0.007+ (0.004)
Num.Obs.	674437	43257	631180	494181
Sample	Full	ZLFN	Not ZLFN	Full

Notes: This table reports difference-in-differences estimates of the effect of exposure to clandestine drug labs on school dropout and household labor income. The dependent variable in columns 1-4 is a dummy for dropping out of high school (Panel A) or middle school (Panel B). The dependent variable in columns 5-7 is the log of household labor income per capita (col. 5), per capita household labor income excluding the individual's own income (col. 6), and per capita household labor income considering only income of members aged 25+ (col. 7). Column 1 uses the full sample. Column 2 restricts to municipalities in the Zona Libre de la Frontera Norte (ZLFN). Columns 3-4 exclude ZLFN municipalities, with column 4 additionally controlling for household income per capita. All columns restrict to individuals of high school age (Panel A) or middle school age (Panel B). All specifications include commuting zone, year-quarter, and quarter fixed effects, and control for sex, age, and household schooling of members aged 25+. Observations are weighted using normalized survey weights. Standard errors clustered at the commuting zone level are reported in parentheses. + $p < 0.1$, * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

[↔ Return to text](#)

Table 14: Heterogeneity by COVID-19 severity

	Interaction with		Excluding CZ with top			
	(1)	(2)	Cases		Deaths	
			(3)	(4)	(5)	(6)
<i>Panel A: Middle school dropout</i>						
Exposure $\times$ Post	-0.015** (0.005)	-0.013* (0.005)	-0.016+ (0.010)	-0.017* (0.008)	-0.011 (0.008)	-0.015 (0.010)
Exposure $\times$ Post $\times$ Cases p.c.	0.001* (0.000)					
Exposure $\times$ Post $\times$ Deaths p.c.		0.004 (0.003)				
N	665,100	665,100	171,289	258,508	278,045	164,793
Panel B: High school dropout						
Exposure $\times$ Post	-0.024** (0.008)	-0.024** (0.008)	-0.046* (0.019)	-0.044*** (0.012)	-0.039** (0.015)	-0.053** (0.018)
Exposure $\times$ Post $\times$ Cases p.c.	0.001* (0.001)					
Exposure $\times$ Post $\times$ Deaths p.c.		0.009** (0.003)				
N	674,437	674,437	168,072	255,358	276,206	160,038

Notes: This table reports difference-in-differences estimates of the effect of increased synthetic drug production on school dropout. The dependent variable is an indicator for not being enrolled in school among individuals of the relevant school age: ages 12-14 for middle school (Panel A) and ages 15-17 for high school (Panel B). Exposure is an indicator for any methamphetamine laboratory seized in the commuting zone before 2015. Post is an indicator for periods after April 2019. Columns (1)-(2) interact the treatment effect with COVID-19 cases and deaths per 1,000 population in 2020, respectively. Columns (3)-(4) exclude commuting zones in the top 5% and 10% of COVID-19 cases. Columns (5)-(6) exclude commuting zones in the top 5% and 10% of COVID-19 deaths. All specifications include commuting zone and year fixed effects and control for gender, age, and household education. Standard errors clustered at the commuting zone level are reported in parentheses. + $p < 0.1$, * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

[↔ Return to text](#)

Table 15: Robustness to Prospera rollback

term	Dependent variable: School dropout			
	Interaction with Prospera		Excluding rural CZ	
	(1)	(2)	(3)	(4)
Panel A: Middle school age				
Exposure $\times$ Post	-0.012* (0.005)	-0.005 (0.005)	-0.014** (0.005)	-0.013* (0.005)
Post $\times$ Prospera Top 25%	0.045** (0.016)			
Post $\times$ Prospera Top 50%		0.055*** (0.011)		
Exposure $\times$ Post $\times$ Prospera Top 25%	-0.132*** (0.025)			
Exposure $\times$ Post $\times$ Prospera Top 50%		-0.063** (0.024)		
Num.Obs.	665100	665100	654639	632029
Panel B: High school age				
Exposure $\times$ Post	-0.019* (0.008)	-0.014+ (0.007)	-0.022*** (0.008)	-0.019* (0.008)
Post $\times$ Prospera Top 25%	0.073*** (0.020)			
Post $\times$ Prospera Top 50%		0.056*** (0.015)		
Exposure $\times$ Post $\times$ Prospera Top 25%	-0.238*** (0.029)			
Exposure $\times$ Post $\times$ Prospera Top 50%		-0.082* (0.040)		
Num.Obs.	674437	674437	665071	643988

Notes: This table reports difference-in-differences estimates of the effect of increased synthetic drug production on school dropout, testing robustness to the rollback of the Prospera conditional cash transfer program. The dependent variable is an indicator for not being enrolled in school among individuals of the relevant school age: ages 1214 for middle school (Panel A) and ages 1517 for high school (Panel B). Exposure is an indicator for any methamphetamine laboratory seized in the commuting zone before 2015. Post is an indicator for periods after April 2019. Columns (1)(2) interact the treatment effect with indicators for high Prospera penetration, defined as commuting zones in the top 25th percentile (column 1) or top 50th percentile (column 2) of the 2017 beneficiary share distribution (ratio of enrolled households to total households from the 2010 census). Columns (3)(4) exclude the most rural commuting zones, where Prospera penetration was highest. All specifications include commuting zone and year fixed effects and control for gender, age, and household education. Standard errors clustered at the commuting zone level are reported in parentheses + $p < 0.1$, * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

[↔ Return to text](#)

G.2 Survey non-response

Table G.2 reports difference-in-differences estimates of exposure on survey non-response. Panel A is unit non-response at the dwelling $\times$ year level: the share of inhabited dwellings that are outright refusals, refusals plus “other reason,” and any non-complete visit. Panel B is item non-response among employed respondents: missing income, income heaping at round values, and missing occupation. Exposure raises outright refusals but leaves item non-response essentially unchanged, indicating that whole households or workers opt out rather than respondents concealing a single attribute.

Table G.2: Exposure to synthetic drug production and survey non-response

Panel A. Unit non-response (dwelling $\times$ year)				
	Refusals only	Refusals + Otro motivo	Any non-complete	
Exposure $\times$ Post	0.007*	0.012	0.018	
	(0.003)	(0.012)	(0.013)	
N	5,489,535	5,489,535	5,489,535	
Control mean	0.027	0.052	0.078	

Panel B. Item non-response among employed respondents				
	Income missing	Heap at 1,000	Heap at 5,000	Occupation missing
Exposure $\times$ Post	0.015	-0.009+	-0.003	-0.001
	(0.014)	(0.005)	(0.003)	(0.001)
N	7,554,458	5,574,481	5,574,481	7,893,739
Control mean	0.258	0.205	0.054	0.001

Notes: Difference-in-differences estimates; standard errors clustered by commuting zone in parentheses. Both panels are at the year level with commuting-zone, year, and quarter fixed effects and commuting-zone linear trends. *Exposure* equals one for commuting zones with a methamphetamine laboratory seized before 2015; *Post* equals one from 2019 (China’s fentanyl ban). “Control mean” is the mean of the dependent variable in the untreated (*Exposure* = 0 or pre-ban) sample.

Panel A is estimated on all inhabited dwellings and is unweighted, since ENOE expansion factors are zero for non-respondents by construction. Each outcome is a 0/1 indicator for the dwelling’s disposition: *Refusals only* equals one if the occupants refused the interview; *Refusals + Otro motivo* adds dwellings not interviewed for “another reason (dwelling inhabited)”; *Any non-complete* equals one for any inhabited dwelling that did not yield a completed interview.

Panel B is estimated on employed respondents, weighted by survey weights, and controls for sex, age, and household adult schooling. *Income missing* is the share of paid workers reporting zero or missing labor income. *Heap at 1,000* and *Heap at 5,000* are the share of *positive* reported incomes that are exact multiples of 1,000 and 5,000 pesos—a proxy for rounded, imprecise income reporting (higher values indicate less precise reports). *Occupation missing* is the share of workers with an unspecified occupation (SINCO code 9999). + $p < 0.1$, * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

[↔ Return to text](#)

References

Daron Acemoglu and David Autor. Chapter 12 - Skills, Tasks and Technologies: Implications for Employment and Earnings. volume 4 of *Handbook of labor economics*, pages 1043–1171. Else-

- vier, 2011. [https://doi.org/10.1016/S0169-7218\(11\)02410-5](https://doi.org/10.1016/S0169-7218(11)02410-5). URL <https://www.sciencedirect.com/science/article/pii/S0169721811024105>. ISSN: 1573-4463.
- Anjali Adukia, Sam Asher, and Paul Novosad. Educational Investment Responses to Economic Opportunity: Evidence from Indian Road Construction. *American Economic Journal: Applied Economics*, 12(1):348–376, January 2020. ISSN 1945-7782. 10.1257/app.20180036. URL <https://www.aeaweb.org/articles?id=10.1257/app.20180036>.
- Lorenzo Aldeco Leo, Andrés Jurado, and Aurora A. Ramírez-Álvarez. Internal migration and drug violence in Mexico. *Journal of Development Economics*, 171:103334, October 2024. ISSN 0304-3878. 10.1016/j.jdeveco.2024.103334. URL <https://www.sciencedirect.com/science/article/pii/S030438782400083X>.
- Alberto Alesina, Salvatore Piccolo, and Paolo Pinotti. Organized Crime, Violence, and Politics. *The Review of Economic Studies*, 86(2):457–499, March 2019. ISSN 0034-6527. 10.1093/restud/rdy036. URL <https://doi.org/10.1093/restud/rdy036>.
- Abby Alpert, David Powell, and Rosalie Liccardo Pacula. Supply-Side Drug Policy in the Presence of Substitutes: Evidence from the Introduction of Abuse-Deterrent Opioids. *American Economic Journal: Economic Policy*, 10(4):1–35, 2018. 10.1257/pol.20170082.
- Parker Asmann. Methamphetamine Taking Over Mexico’s Domestic Drug Market, April 2021. URL <https://insightcrime.org/news/methamphetamine-mexico-drug-market/>.
- Susan Athey and Guido W Imbens. The state of applied econometrics: Causality and policy evaluation. *Journal of Economic Perspectives*, 31(2):3–32, 2017.
- Susan Athey and Guido W Imbens. Machine learning methods that economists should know about. *Annual Review of Economics*, 11:685–725, 2019.
- David Atkin. Endogenous Skill Acquisition and Export Manufacturing in Mexico. *American Economic Review*, 106(8):2046–2085, August 2016. ISSN 0002-8282. 10.1257/aer.20120901. URL <https://www.aeaweb.org/articles?id=10.1257/aer.20120901>.
- David H. Autor, Frank Levy, and Richard J. Murnane. The Skill Content of Recent Technological Change: An Empirical Exploration*. *The Quarterly Journal of Economics*, 118(4):1279–1333, November 2003. ISSN 0033-5533. 10.1162/003355303322552801. URL <https://doi.org/10.1162/003355303322552801>.
- Marco Battaglini, Luigi Guiso, Chiara Lacava, Douglas L. Miller, and Eleonora Patacchini. Refining public policies with machine learning: The case of tax auditing. *Journal of Econometrics*, 249:105847, May 2025. ISSN 0304-4076. 10.1016/j.jeconom.2024.105847. URL <https://www.sciencedirect.com/science/article/pii/S0304407624001921>.

- Maria Laura Battagliola and others. Refining public policies with machine learning: The case of tax auditing. *Journal of Econometrics*, 2024.
- Gary S. Becker. Crime and Punishment: An Economic Approach. *Journal of Political Economy*, 76(2):169–217, 1968. ISSN 0022-3808. URL <https://www.jstor.org/stable/1830482>. Publisher: University of Chicago Press.
- Brian Bell, Rui Costa, and Stephen Machin. Why Does Education Reduce Crime? *Journal of Political Economy*, 130(3):732–765, March 2022. ISSN 0022-3808. 10.1086/717895. URL <https://www.journals.uchicago.edu/doi/abs/10.1086/717895>. Publisher: The University of Chicago Press.
- Christopher Blattman, Gustavo Duncan, Benjamin Lessing, and Santiago Tobón. Gang rule: Understanding and countering criminal governance. *The Review of Economic Studies*, 92(3):1497–1531, 2025.
- Ryan Brown and Andrea Velásquez. The effect of violent crime on the human capital accumulation of young adults. *Journal of Development Economics*, 127:1–12, 2017.
- Petri Böckerman and Mika Haapanen. The effect of polytechnic reform on migration. *Journal of Population Economics*, 26(2):593–617, 2013. URL <https://ideas.repec.org//a/spr/jopoec/v26y2013i2p593-617.html>. Publisher: Springer & European Society for Population Economics.
- Mariana Calderón, Josué Cortés, Jorge Pérez Pérez, and Alejandrina Salcedo. Disentangling the Effects of Large Minimum Wage and VAT Changes on Prices: Evidence from Mexico. *Labour Economics*, 80:102294, January 2023. ISSN 0927-5371. 10.1016/j.labeco.2022.102294. URL <https://www.sciencedirect.com/science/article/pii/S0927537122001841>.
- Nadia Campaniello, Rowena Gray, and Giovanni Mastrobuoni. Returns to education in criminal organizations: Did going to college help Michael Corleone? *Economics of Education Review*, 54: 242–258, 2016.
- Gian Maria Campedelli, Gianmarco Daniele, and Marco Le Moglie. Mafia, politics and machine predictions. CEPR Discussion Paper DP19322, Centre for Economic Policy Research, 2024.
- Javier Cano-Urbina. The role of the informal sector in the early careers of less-educated workers. *Journal of Development Economics*, 112:33–55, January 2015. ISSN 0304-3878. 10.1016/j.jdeveco.2014.10.002. URL <https://www.sciencedirect.com/science/article/pii/S0304387814001084>.
- Juan Camilo Castillo and Dorothy Kronick. The Logic of Violence in Drug War. *American Political Science Review*, 114(3):874–887, August 2020. ISSN 0003-0554, 1537-5943. 10.1017/S0003055420000246. URL <https://www.cambridge.org/core/journals/american-political-science-review/article/abs/logic-of-violence-in-drug-war/A20A39930D2273FFA3E7750F5A5A66E5>.

- Kerwin Kofi Charles, Erik Hurst, and Matthew J. Notowidigdo. Housing Booms and Busts, Labor Market Opportunities, and College Attendance. *American Economic Review*, 108(10):2947–2994, October 2018. ISSN 0002-8282. 10.1257/aer.20151604. URL <https://www.aeaweb.org/articles?id=10.1257/aer.20151604>.
- Laura Chioda, João M. P. De Mello, and Rodrigo R. Soares. Spillovers from conditional cash transfer programs: Bolsa Família and crime in urban Brazil. *Economics of Education Review*, 54:306–320, 2016.
- Congressional Research Service. Illicit Fentanyl and China’s Role. CRS In Focus IF10890, Congressional Research Service, 2019. URL https://www.congress.gov/crs_external_products/IF/HTML/IF10890.web.html.
- Gianmarco Daniele, Marco Le Moglie, and Federico Masera. Pains, guns and moves: The effect of the U.S. opioid epidemic on Mexican migration. *Journal of Development Economics*, 160: 102983, January 2023. ISSN 0304-3878. 10.1016/j.jdeveco.2022.102983. URL <https://www.sciencedirect.com/science/article/pii/S0304387822001250>.
- Daniel de Roux, Boris Perez, Andrés Moreno, Maria del Pilar Villamil, and César Figueroa. Tax Fraud Detection for Under-Reporting Declarations Using an Unsupervised Machine Learning Approach. In *Proceedings of the 24th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*, KDD ’18, pages 215–222, New York, NY, USA, July 2018. Association for Computing Machinery. ISBN 978-1-4503-5552-0. 10.1145/3219819.3219878. URL <https://doi.org/10.1145/3219819.3219878>.
- Secretaría de Salud. Dirección General de Epidemiología - Datos Abiertos Bases Históricas, 2025. URL <http://www.gob.mx/salud/documentos/datos-abiertos-bases-historicas-direccion-general-de-epidemiologia>.
- Adeline Delavande and Basit Zafar. University Choice: The Role of Expected Earnings, Nonpecuniary Outcomes, and Financial Constraints. *Journal of Political Economy*, 127(5):2343–2393, October 2019. ISSN 0022-3808. 10.1086/701808. URL <https://www.journals.uchicago.edu/doi/abs/10.1086/701808>. Publisher: The University of Chicago Press.
- Melissa Dell. Trafficking Networks and the Mexican Drug War. *American Economic Review*, 105(6): 1738–1779, June 2015. ISSN 0002-8282. 10.1257/aer.20121637. URL <https://www.aeaweb.org/articles?id=10.1257/aer.20121637>.
- Rafael Dix-Carneiro, Rodrigo R. Soares, and Gabriel Ulyssea. Economic shocks and crime: Evidence from the Brazilian trade liberalization. *American Economic Journal: Applied Economics*, 10(4): 158–195, 2018. 10.1257/app.20170080.
- Oeindrila Dube and Juan F. Vargas. Commodity price shocks and civil conflict: Evidence from Colombia. *The Review of Economic Studies*, 80(4):1384–1421, 2013. 10.1093/restud/rdt009.

- Steven Dudley. El flujo de precursores químicos para la producción de drogas sintéticas en México. Technical report, InSight Crime, 2023.
- Gilles Duranton. A Proposal to Delineate Metropolitan Areas in Colombia. *Revista Desarrollo y Sociedad*, (75):223–264, January 2015. ISSN 1900-7760. 10.13043/dys.75.6. URL <https://revistas.uniandes.edu.co/index.php/dys/article/view/6673>. Number: 75.
- Isaac Ehrlich. Participation in Illegitimate Activities: A Theoretical and Empirical Investigation. *Journal of Political Economy*, 81(3):521–565, 1973. ISSN 0022-3808. URL <https://www.jstor.org/stable/1831025>. Publisher: The University of Chicago Press.
- José Ramón Enríquez, Horacio Larreguy, and Alberto Simpser. The impact of confounders, spillovers and interactions on social distancing policy effects estimates. *Scientific Reports*, 15(1):34350, October 2025. ISSN 2045-2322. 10.1038/s41598-025-16861-1. URL <https://www.nature.com/articles/s41598-025-16861-1>. Publisher: Nature Publishing Group.
- William N. Evans, Ethan M. J. Lieber, and Patrick Power. How the Reformulation of OxyContin Ignited the Heroin Epidemic. *The Review of Economics and Statistics*, 101(1):1–15, 2019.
- Joseph Friedman and Chelsea L. Shover. Charting the fourth wave: Geographic, temporal, race/ethnicity and demographic trends in polysubstance fentanyl overdose deaths in the United States, 2010–2021. *Addiction*, 118(12):2477–2485, 2023. 10.1111/add.16318.
- Sara Garcia. Químicos y cocineros: así mata el fentanilo, May 2023. URL <https://insightcrime.org/es/investigaciones/quimicos-cocineros-fentanilo-mata/>.
- Daniel Haanwinckel. Does regional variation in wage levels identify the effects of a national minimum wage?, April 2025. URL <http://arxiv.org/abs/2307.01284>. arXiv:2307.01284 [econ].
- International Narcotics Control Board. Precursors and chemicals frequently used in the illicit manufacture of narcotic drugs and psychotropic substances: Report of the international narcotics control board for 2022. Technical Report E/INCB/2022/4, United Nations, Vienna, 2023. URL https://www.incb.org/documents/Publications/AnnualReports/AR2022/Precursors_Report/E_INCB_2022_4_eng.pdf.
- Brenda Jarillo, Beatriz Magaloni, Edgar Franco, and Gustavo Robles. How the Mexican Drug War affects kids and schools? Evidence on effects and mechanisms. *International Journal of Educational Development*, 51:135–146, 2016.
- Lawrence F. Katz and Kevin M. Murphy. Changes in Relative Wages, 1963–1987: Supply and Demand Factors*. *The Quarterly Journal of Economics*, 107(1):35–78, February 1992. ISSN 0033-5533. 10.2307/2118323. URL <https://doi.org/10.2307/2118323>.
- Jon Kleinberg, Jens Ludwig, Sendhil Mullainathan, and Ziad Obermeyer. Prediction policy problems. *American Economic Review*, 105(5):491–495, 2015.

- Jon Kleinberg, Himabindu Lakkaraju, Jure Leskovec, Jens Ludwig, and Sendhil Mullainathan. Human Decisions and Machine Predictions*. *The Quarterly Journal of Economics*, 133(1):237–293, February 2018. ISSN 0033-5533. 10.1093/qje/qjx032. URL <https://doi.org/10.1093/qje/qjx032>.
- Quinto Elemento Lab. Narco Contaminación: La catástrofe ambiental de los laboratorios clandestinos, 2025. URL <https://quintoelab.org/narco-contaminacion/catastrofe-ambiental-laboratorios-clandestinos/>.
- Lance Lochner and Enrico Moretti. The Effect of Education on Crime: Evidence from Prison Inmates, Arrests, and Self-Reports. *American Economic Review*, 94(1):155–189, March 2004. ISSN 0002-8282. 10.1257/000282804322970751. URL <https://www.aeaweb.org/articles?id=10.1257/000282804322970751>.
- Stephen Machin and Matteo Sandi. Crime and education. *Annual Review of Economics*, 17(1):241–260, August 2025. ISSN 1941-1383. 10.1146/annurev-economics-090324-035606. URL <https://researchonline.lse.ac.uk/id/eprint/129152/>. Num Pages: 20.
- Stephen Machin, Olivier Marie, and Sunica Vuji. The Crime Reducing Effect of Education. *The Economic Journal*, 121(552):463–484, 2011. ISSN 1468-0297. 10.1111/j.1468-0297.2011.02430.x. URL <https://onlinelibrary.wiley.com/doi/abs/10.1111/j.1468-0297.2011.02430.x>. _eprint: <https://onlinelibrary.wiley.com/doi/pdf/10.1111/j.1468-0297.2011.02430.x>.
- Fernanda Marquez-Padilla, Susan Parker, and Tom Vogl. Rolling Back Progres: School and Work After the End of a Landmark Anti-Poverty Program. Technical Report w33527, National Bureau of Economic Research, Cambridge, MA, February 2025. URL <http://www.nber.org/papers/w33527.pdf>.
- Daniel Mejía and Pascual Restrepo. The economics of the war on illegal drug production and trafficking. *Journal of Economic Behavior & Organization*, 126:255–275, 2016.
- Litterio Mirenda, Sauro Mocetti, and Lucia Rizzica. The Economic Effects of Mafia: Firm Level Evidence. *American Economic Review*, 2022. ISSN 0002-8282. 10.1257/aer.20201015. URL <https://www.aeaweb.org/articles?id=10.1257/aer.20201015&from=f>.
- Sendhil Mullainathan and Jann Spiess. Machine learning: An applied econometric approach. *Journal of Economic Perspectives*, 31(2):87–106, 2017.
- National Drug Intelligence Center. Gulf Coast High Intensity Drug Trafficking Area Drug Market Analysis 2009. Drug market analysis, U.S. Department of Justice, National Drug Intelligence Center, 2009. URL <https://www.justice.gov/archive/ndic/pubs32/32769/production.htm>.
- National Institute on Drug Abuse. Fentanyl DrugFacts, 2021. URL <https://nida.nih.gov/publications/drugfacts/fentanyl>.

- Susan W. Parker and Petra E. Todd. Conditional Cash Transfers: The Case of Progresa/Oportunidades. *Journal of Economic Literature*, 55(3):866–915, September 2017. ISSN 0022-0515. 10.1257/jel.20151233. URL <https://www.aeaweb.org/articles?id=10.1257/jel.20151233>.
- Paolo Pinotti. The Economic Costs of Organised Crime: Evidence from Southern Italy. *The Economic Journal*, 125(586):F203–F232, 2015. ISSN 0013-0133. URL <https://www.jstor.org/stable/24737566>. Publisher: [Royal Economic Society, Wiley].
- Rafael Prieto-Curiel, Gian Maria Campedelli, and Alejandro Hope. Reducing cartel recruitment is the only way to lower violence in Mexico. *Science*, 381(6664):1312–1316, 2023. 10.1126/science.adh2888.
- Sam Quinones. *The Least of Us: True Tales of America and Hope in the Time of Fentanyl and Meth*. Bloomsbury Publishing, New York, 2021.
- Manisha Shah and Bryce Millett Steinberg. Drought of Opportunities: Contemporaneous and Long-Term Impacts of Rainfall Shocks on Human Capital. *Journal of Political Economy*, 125(2):527–561, April 2017. ISSN 0022-3808. 10.1086/690828. URL <https://www.journals.uchicago.edu/doi/abs/10.1086/690828>. Publisher: The University of Chicago Press.
- J. M. C. Santos Silva and Silvana Tenreyro. The Log of Gravity. *The Review of Economics and Statistics*, 88(4):641–658, November 2006. ISSN 0034-6535. 10.1162/rest.88.4.641. URL <https://doi.org/10.1162/rest.88.4.641>.
- Fernanda Sobrino. Mexican cartel wars: Fighting for the US opioid market. 2019.
- Maria Micaela Sviatschi. Making a NARCO: Childhood Exposure to Illegal Labor Markets and Criminal Life Paths. *Econometrica*, 90(4):1835–1878, 2022. ISSN 1468-0262. 10.3982/ECTA17082. URL <https://onlinelibrary.wiley.com/doi/abs/10.3982/ECTA17082>. _eprint: <https://onlinelibrary.wiley.com/doi/pdf/10.3982/ECTA17082>.
- United Nations Office on Drugs and Crime. Three precursors of the most common synthesis routes used in illicit fentanyl manufacture now under international control. UNODC Laboratory and Scientific Service, Early Warning Advisory, April 2022. URL <https://www.unodc.org/LSS/Announcement/Details/b152bda5-5d71-4f7e-9d68-1bdd9af04a83>.
- United Nations Office on Drugs and Crime. Crime-related illicit financial flows: latest progress. Technical report, United Nations Office on Drugs and Crime (UNODC), Vienna, 2023. URL https://www.unodc.org/documents/data-and-analysis/IFF/2023/IFFs_Estimates_Report_2023-final-11dec2023.pdf.
- U.S. Department of the Treasury. Treasury Sanctions Illicit Fentanyl Supply Network Supporting the Sinaloa Cartel, October 2025. URL <https://home.treasury.gov/news/press-releases/sb0272>. Press Release No. SB0272ç.

- U.S. Drug Enforcement Administration. Fentanyl Flow in the United States. Intelligence Report DEA-DCT-DIR-008-20, U.S. Drug Enforcement Administration, 2020. URL https://www.dea.gov/sites/default/files/2020-03/DEA_GOV_DIR-008-20%20Fentanyl%20Flow%20in%20the%20United%20States_0.pdf.
- U.S. Drug Enforcement Administration. DEA Laboratory Testing Reveals that 6 out of 10 Fentanyl-Laced Fake Prescription Pills Now Contain a Potentially Lethal Dose of Fentanyl, 2022. URL <https://www.dea.gov/alert/dea-laboratory-testing-reveals-6-out-10-fentanyl-laced-fake-prescription-pills-now-contain>.
- U.S. Immigration and Customs Enforcement. United states–mexico bi-national criminal proceeds study. Cornerstone report, U.S. Immigration and Customs Enforcement (ICE), 2010. URL <https://www.ice.gov/doclib/cornerstone/pdf/cps-study.pdf>.
- U.S. Senate Permanent Subcommittee on Investigations. Combatting the Opioid Crisis: Exploiting Vulnerabilities in International Mail. Staff report, Committee on Homeland Security and Governmental Affairs, United States Senate, 2018. URL [https://www.hsgac.senate.gov/wp-content/uploads/imo/media/doc/Fentanyl%20Synthetic%20Drugs%20Report%2063%20\(FINAL\)%20-%203.pdf](https://www.hsgac.senate.gov/wp-content/uploads/imo/media/doc/Fentanyl%20Synthetic%20Drugs%20Report%2063%20(FINAL)%20-%203.pdf).
- Art Van Zee. The Promotion and Marketing of OxyContin: Commercial Triumph, Public Health Tragedy. *American Journal of Public Health*, 99(2):221–227, 2009. 10.2105/AJPH.2007.131714.